

The Ultimate PyTorch Technical Reference

Deep Learning Notes

February 2026

Contents

1	Extensive Tensor Operations and "Weird" Logic	2
1.1	Standard and Structural Initializers	2
1.2	Advanced Indexing and Selection	2
1.3	Shape Manipulation and Memory Management	2
2	Comprehensive Loss Function Catalog	3
2.1	Deep Dive: When to use what?	3
3	Implementation of the Training Pipeline	3
3.1	The Anatomy of a Training Loop	3
3.2	Critical Checklist	4

1 Extensive Tensor Operations and "Weird" Logic

Tensors are the bedrock of PyTorch. Beyond simple addition, these functions allow for complex data manipulation.

1.1 Standard and Structural Initializers

- `torch.eye(n, m)`: Returns a 2D tensor with ones on the diagonal (Identity matrix).
- `torch.diag(input)`: Extracts a diagonal or constructs a diagonal matrix.
- `torch.tril(input) / torch.triu(input)`: Returns the lower or upper triangular part of a matrix (used for masking in Transformers).
- `torch.full((shape), value)`: Creates a tensor filled entirely with a specific scalar.
- `torch.linspace(start, end, steps)`: Creates a one-dimensional tensor of equally spaced points.

1.2 Advanced Indexing and Selection

These functions are the "weird" workhorses for dynamic neural architectures.

- `torch.einsum(equation, *operands)`: Einstein summation. Handles products, sums, and transposes in one string.
 - Batch MatMul: `'bij,bjk->bik'`
 - Inner product: `'i,i->'`
 - Trace: `'ii'`
- `torch.gather(input, dim, index)`: Gathers values along an axis.
- `torch.scatter_(dim, index, src)`: Writes values into a tensor at specific indices (in-place).
- `torch.where(condition, x, y)`: Returns a tensor of elements selected from either x or y based on the condition.
- `torch.roll(input, shifts, dims)`: Circularly shifts the tensor along given dimensions.
- `torch.unbind(input, dim)`: Removes a dimension, returning a tuple of all slices along that dimension.

1.3 Shape Manipulation and Memory Management

- `tensor.view(*shape)`: Changes shape without copying data (must be contiguous).
- `tensor.reshape(*shape)`: Changes shape, may copy data if necessary.
- `tensor.permute(*dims)`: Arbitrary reordering of dimensions.
- `tensor.squeeze() / tensor.unsqueeze(dim)`: Removes or adds dimensions of size 1.
- `tensor.expand(*sizes)`: Views a singleton dimension as larger without using extra memory.
- `tensor.repeat(*sizes)`: Copies the data to repeat the tensor.

2 Comprehensive Loss Function Catalog

Loss functions (`torch.nn`) define the objective the network tries to minimize.

Loss Function	Typical Task	Output Requirement
<code>nn.MSELoss</code>	Regression	Linear/Raw output
<code>nn.L1Loss</code>	Regression (MAE)	Robust to outliers
<code>nn.CrossEntropyLoss</code>	Multi-class	Raw Logits (No Softmax!)
<code>nn.BCEWithLogitsLoss</code>	Binary / Multi-label	Raw Logits (No Sigmoid!)
<code>nn.NLLLoss</code>	Multi-class	LogSoftmax outputs
<code>nn.KLDivLoss</code>	Distillation	Log-probabilities
<code>nn.HuberLoss</code>	Robust Regression	Combines MSE and L1

2.1 Deep Dive: When to use what?

1. **`nn.CrossEntropyLoss`**: Use for classification where each input belongs to exactly one class. It expects raw scores (logits) because it internally applies `LogSoftmax`.
2. **`nn.BCEWithLogitsLoss`**: Use for binary classification (Yes/No) or multi-label classification (e.g., tags in an image). It is more numerically stable than `Sigmoid + BCELoss`.
3. **`nn.MSELoss`**: Standard for predicting continuous values. If your data has many outliers, consider `nn.L1Loss` or `nn.HuberLoss`.

3 Implementation of the Training Pipeline

A training loop in PyTorch is explicit, giving the developer full control over the gradient flow.

3.1 The Anatomy of a Training Loop

```
1 import torch
2 import torch.nn as nn
3 import torch.optim as optim
4
5 # 1. Architecture Definition
6 class ResNetBlock(nn.Module):
7     def __init__(self, in_dims, out_dims):
8         super().__init__()
9         self.linear = nn.Linear(in_dims, out_dims)
10        self.relu = nn.ReLU()
11
12    def forward(self, x):
13        return self.relu(self.linear(x))
14
15 # 2. Setup
16 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
17 model = ResNetBlock(784, 10).to(device)
18 optimizer = optim.Adam(model.parameters(), lr=1e-3)
19 criterion = nn.CrossEntropyLoss()
20
21 # 3. Training Loop
```

```

22 def train_one_epoch(epoch_index):
23     model.train() # Enable dropout/batchnorm
24     running_loss = 0.
25
26     for i, (inputs, labels) in enumerate(train_loader):
27         inputs, labels = inputs.to(device), labels.to(device)
28
29         # MANDATORY: Clear gradients from previous step
30         optimizer.zero_grad()
31
32         # Forward pass
33         outputs = model(inputs)
34         loss = criterion(outputs, labels)
35
36         # Backward pass (Autograd magic)
37         loss.backward()
38
39         # Update weights
40         optimizer.step()
41
42         running_loss += loss.item()
43
44     return running_loss / len(train_loader)
45
46 # 4. Evaluation Loop
47 def validate():
48     model.eval() # Disable dropout/batchnorm
49     total_loss = 0
50     with torch.no_grad(): # Disable gradient calculation (saves memory)
51         for inputs, labels in val_loader:
52             outputs = model(inputs.to(device))
53             loss = criterion(outputs, labels.to(device))
54             total_loss += loss.item()
55     return total_loss / len(val_loader)

```

3.2 Critical Checklist

- `zero_grad()`: Without this, gradients accumulate (sum up) across batches, destroying training.
- `loss.backward()`: Populates the `.grad` attribute for every parameter.
- `optimizer.step()`: Adjusts parameters based on the `.grad` values.
- `torch.no_grad()`: Always use this during inference/validation to prevent the computation graph from being built, which saves a massive amount of VRAM.