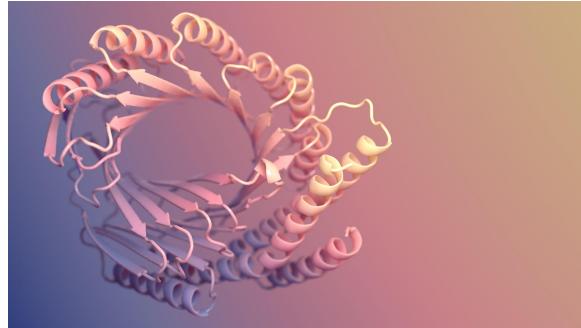


FLOW MATCHING AND DIFFUSION MODELS

EXAMPLES & MOTIVATION



Pi0

RFDiffusion



Stable Diffusion

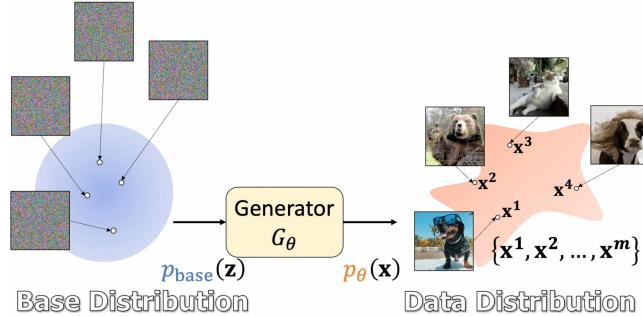
TABLE OF CONTENTS

1. Overview
2. Preliminaries
3. Residual Flows
4. Flow Matching
5. Diffusion
6. Extra Stuff

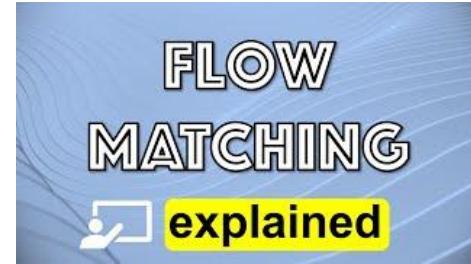
OVERVIEW

GENERAL IDEA

Goal: convert some known probability distribution into a one that describes data well



Imagine deforming substances

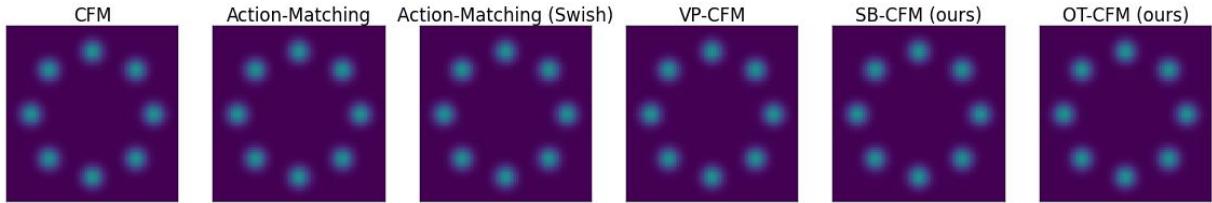


Probability is our substance

HOW DO WE LEARN TO DEFORM DISTRIBUTIONS?

Goal: Convert 8 two-dimensional Gaussians into “two Moons” shaped distribution

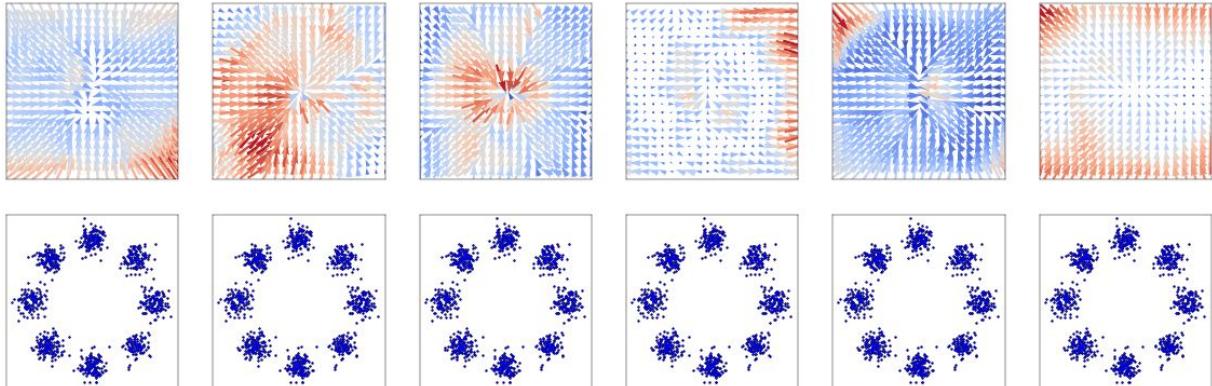
Flow



Underlying
time-dependent
vector field



We want to learn
this vector
field

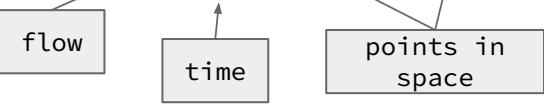


Animations source: [TorchCFM](#)

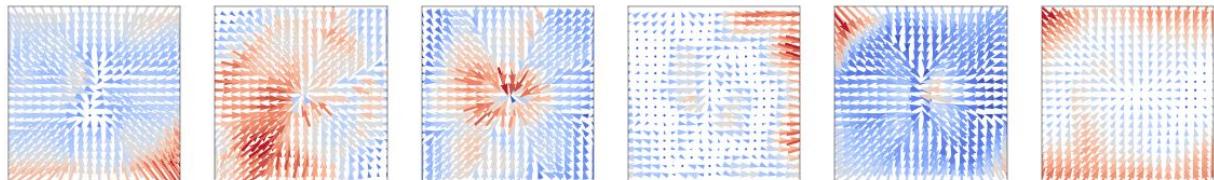
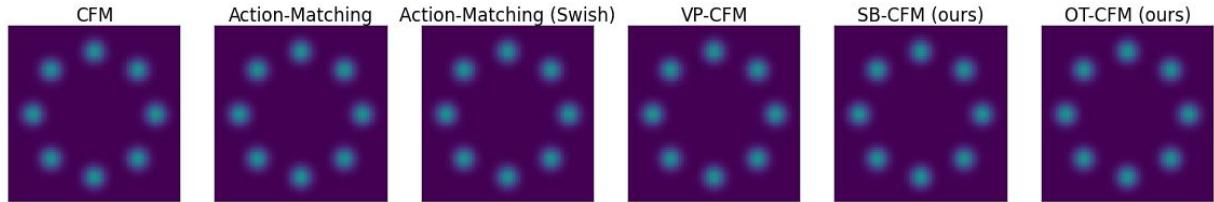
HOW DO WE LEARN TO DEFORM DISTRIBUTIONS?

Goal: Convert 8 two-dimensional Gaussians into “two Moons” shaped distribution

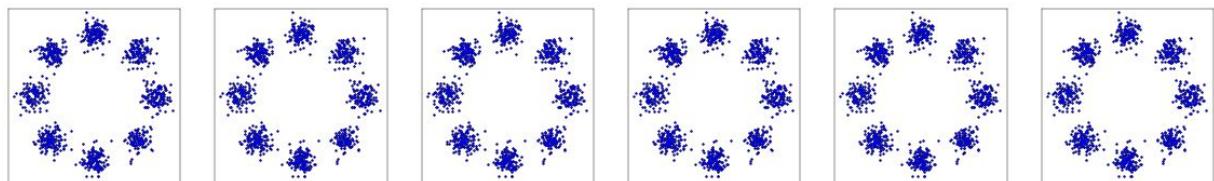
$$\phi : [0, 1] \times \mathbb{R}^d \rightarrow \mathbb{R}^d$$



$$\frac{d\phi_t(x_t)}{dt} = u_t(x_t)$$



$$\mathcal{L}_{FM}(\theta) = \mathbb{E}_{t, p_i(x)} \|v_{\theta, t}(X) - u_t(X)\|$$



Animations source: [TorchCFM](#)

WHY DO WE LEARN TO DEFORM DISTRIBUTIONS?

Generation as sampling from the data distribution

Data distribution: Distribution of objects that we want to generate:

Probability density:

$$p_{\text{data}} : \mathbb{R}^d \rightarrow \mathbb{R}_{\geq 0}, \\ z \mapsto p_{\text{data}}(z)$$

p_{data}

*Note: We
don't know the
probability
density!*

Generation means sampling the data distribution:

$$z \sim p_{\text{data}}$$



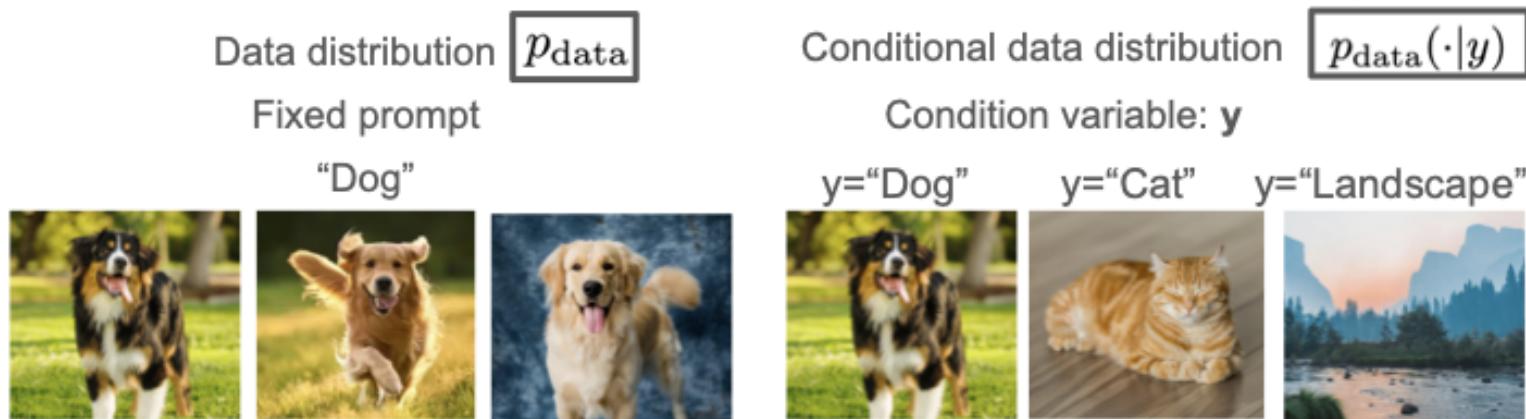
$$z =$$



So in the end we do not necessarily want to know the distribution, but we want to **be able to sample from it**

WHY DO WE LEARN TO DEFORM DISTRIBUTIONS?

Conditional Generation allows us to condition on prompts



Conditional generation means sampling the conditional data distribution:

$$z \sim p_{\text{data}}(\cdot|y)$$

We will first focus on unconditional generation and then learn how to translate an unconditional model to a conditional one.

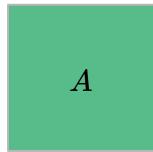
PRELIMINARIES

DEFORMING PROBABILITY

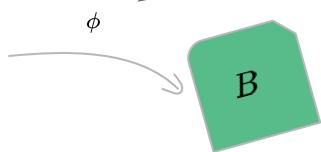
Consider invertible map:

$$Y = \phi(X)$$

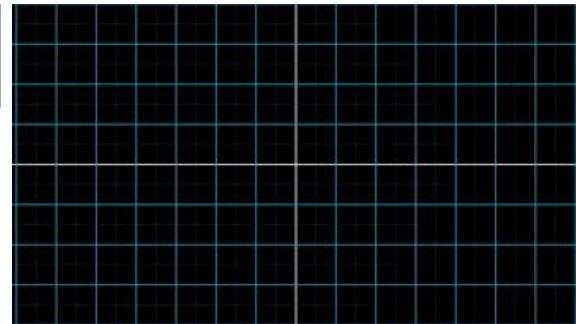
Random variables



The map determines how space is deformed

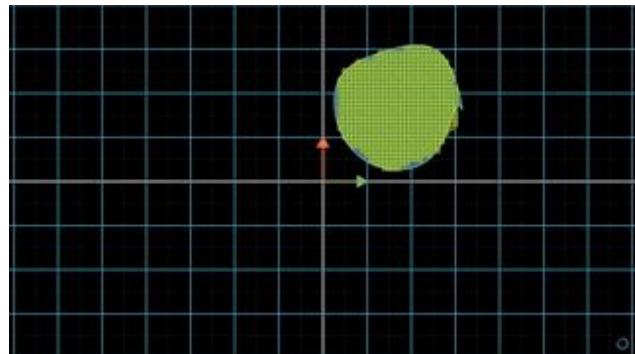


ϕ



Probabilities have to stay the same!

$$\mathbb{P}[Y \in B] = \mathbb{P}[X \in A] \implies \int_B f_Y(y) dy = \int_A f_X(x) dx$$

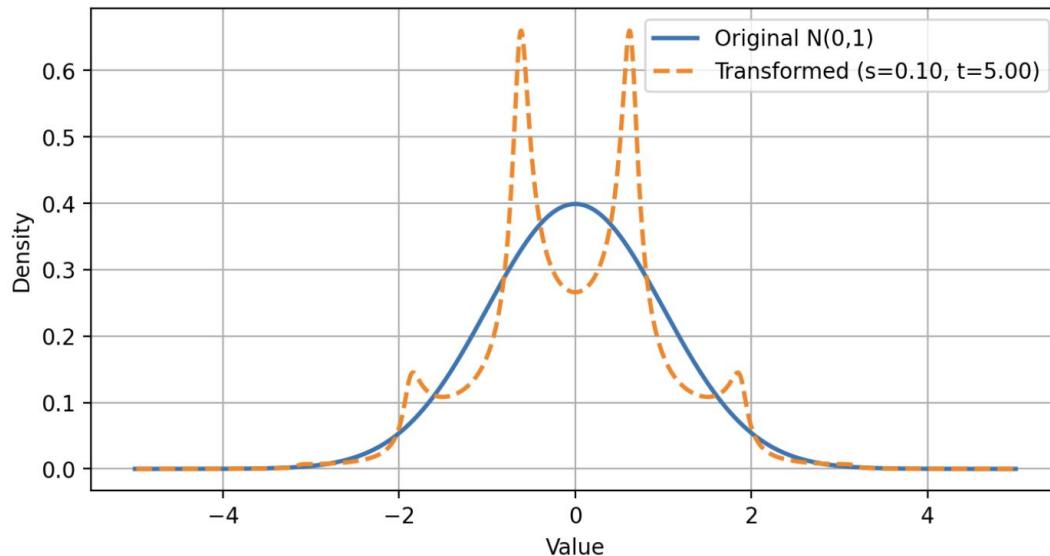


$$\int_A f_X(x) dx = \int_B f_X(\phi^{-1}(y)) |\det D\phi^{-1}(y)| dy$$

$$\implies f_Y(y) = f_X(\phi^{-1}(y)) \cdot |\det D\phi^{-1}(y)|$$

DEFORMING PROBABILITY: EXAMPLE

$X \sim \mathcal{N}(0, 1)$, $Y = \Phi(X) = X + s \sin(tX)$, where $s, t \in \mathbb{R}$ and $|st| < 1$



Managed to transform a unimodal PDF into a multimodal PDF!

How was this possible?

DEFORMING PROBABILITY: GENERAL CASE

$$X \sim p_{\text{init}}, \Phi_{\theta}(X) \sim p_d$$

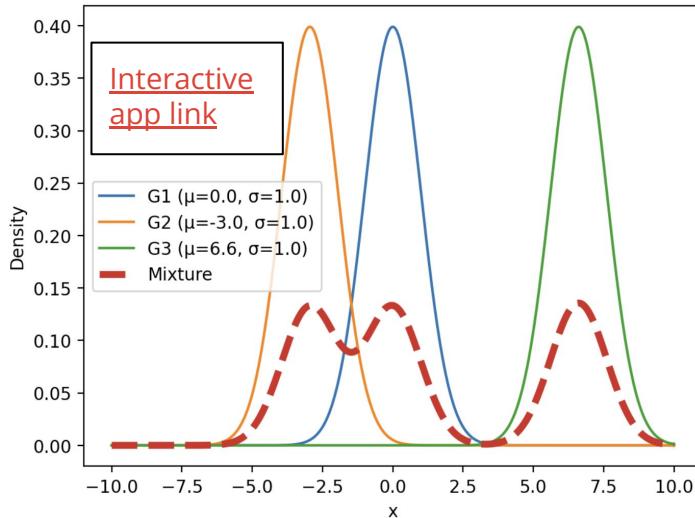
Some known distribution

Deformed distribution

When the deformed distribution will be good?

When it fits the data well!

For example, a Gaussian Mixture (GM) with Gaussians at each observed data point

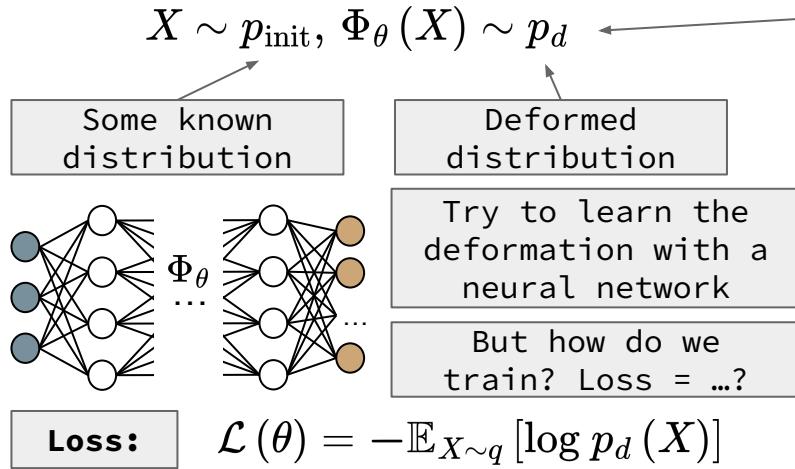


Of course for such a GM one would need at least
#parameters = #data points

So we could control the deformation with parameters. This brings us to...

RESIDUAL FLOWS

LEARNING THE DEFORMATION: KULLBACK-LEIBLER DIVERGENCE



Need to measure how similar deformed distribution is to the true data distribution q

KL divergence:

$$D_{KL}(q, p_d) = \int q(x) \log \frac{q(x)}{p_d(x)}$$

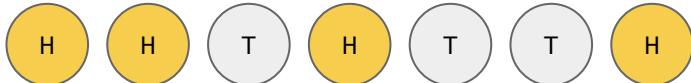
Why???

Consider a coin:

$\begin{cases} q_1 & \text{heads} \\ q_2 & \text{tails} \end{cases}$

$\begin{cases} p_1 & \text{heads} \\ p_2 & \text{tails} \end{cases}$

N obs:



Average likelihood ratio per observation:

$$\begin{aligned} D_{KL}(q, p_d) &= \mathbb{E}_{X \sim q} \left[\log \frac{q(X)}{p_d(X)} \right] \\ &= \mathbb{E}_{X \sim q} [\log q(X)] - \underbrace{\mathbb{E}_{X \sim q} [\log p_d(X)]}_{\text{So called cross-entropy}} \end{aligned}$$

Does not depend on model parameters

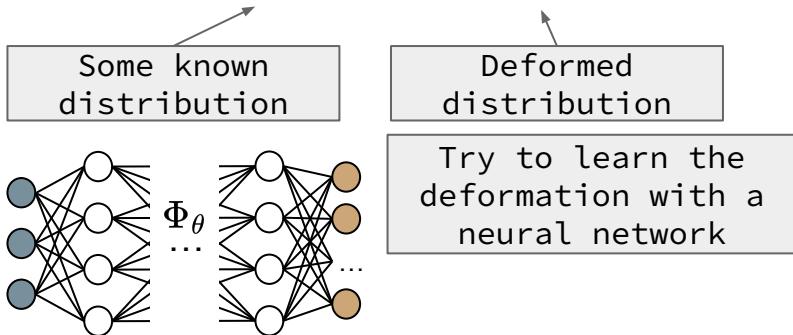
In the limit:

$$\sum_x q(x) \log \frac{q(x)}{p(x)}$$

$$\left(\prod_N \frac{q(x_i)}{p(x_i)} \right)^{\frac{1}{N}} = \left(\frac{q_1^{N_h} q_2^{N_t}}{p_1^{N_h} p_2^{N_t}} \right)^{\frac{1}{N}} \text{ take log: } \frac{N_h}{N} \log \frac{q_1}{p_1} + \frac{N_t}{N} \log \frac{q_2}{p_2}$$

MONTE CARLO ESTIMATES

$$X \sim p_{\text{init}}, \Phi_{\theta}(X) \sim p_d$$



Loss:

$$\mathcal{L}(\theta) = -\mathbb{E}_{X \sim q} [\log p_d(X)]$$

An integral, possibly over an infinite space...

How do we compute such a thing during training?

Approximate:

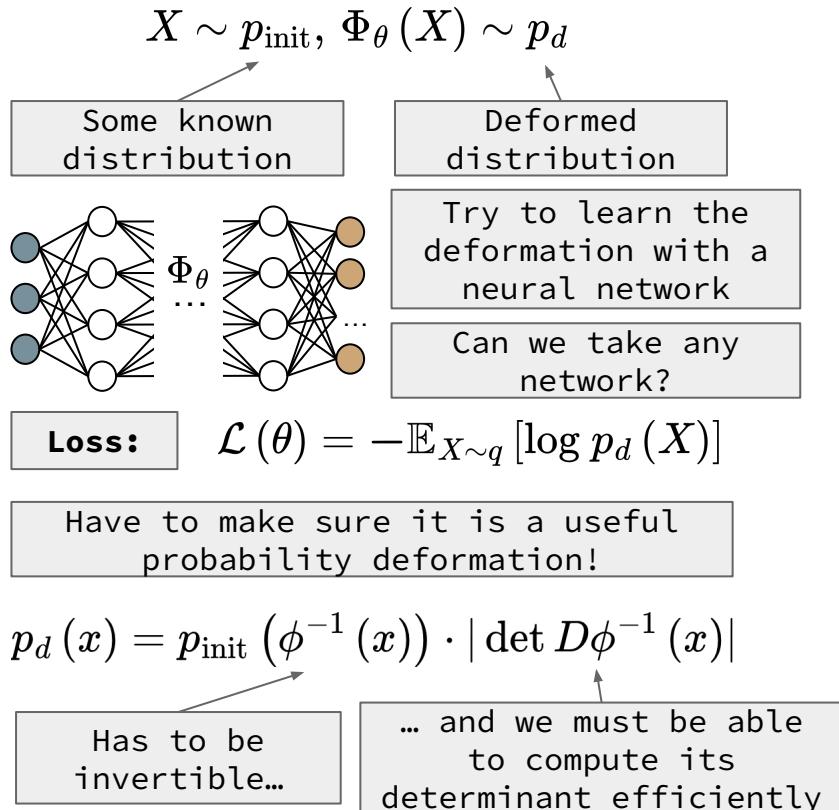
$$\mathbb{E}[f(X)] \approx \frac{1}{N} \sum_i f(x_i)$$

How do we get samples?

DATA!

Samples drawn according to the distribution

LEARNING THE DEFORMATION: RESIDUAL FLOWS



One way to do it - Residual Networks:

$$\phi = \phi_N \circ \phi_{N-1} \circ \cdots \circ \phi_1$$

$$\text{where } \phi_k(x) = x + \delta u_k(x), \delta \in \mathbb{R}$$

Proper u can assure invertibility and easy computation of the determinant (e.g. make the jacobian triangular)

$$\begin{bmatrix} \frac{\partial \phi_1}{\partial x_1} & 0 & \dots & 0 \\ \frac{\partial \phi_2}{\partial x_1} & \frac{\partial \phi_2}{\partial x_2} & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial \phi_d}{\partial x_1} & \frac{\partial \phi_d}{\partial x_2} & \dots & \frac{\partial \phi_d}{\partial x_d} \end{bmatrix}$$

$$\mathcal{L}(\theta) = \mathbb{E}_{X \sim q} \left[\log p_{\text{init}}(\phi^{-1}(x)) + \sum_k \log \det \phi_k \right]$$



RESIDUAL FLOWS AS NUMERICAL INTEGRATION

Looks very similar to Forward Euler integration scheme

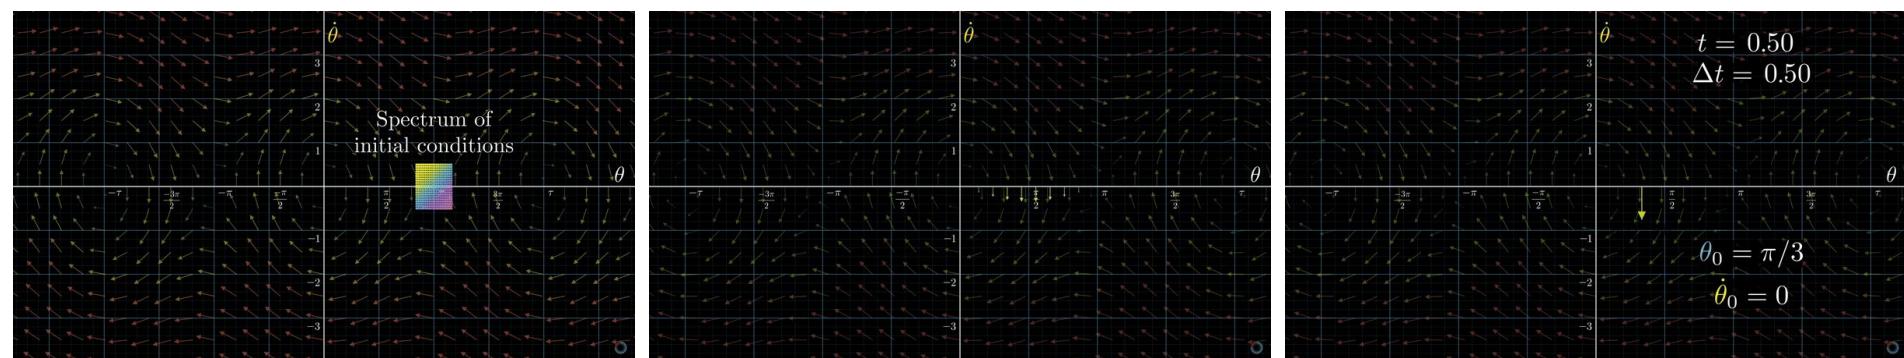
One can notice that this

$$\phi_k(x) = x + \delta \cdot u_k(x), \delta \in \mathbb{R}$$

$$x_{t+\Delta t} \approx x_t + \Delta t \cdot v(x_t, t)$$

They're the same picture.

Numerical integration intuition: we track what's happening with points of a substance in space.



In real life the vector field guides the points along continuous paths

During a computation we use discrete time steps

Changing the timestep affects the result!

The core idea is approximating such Ordinary Differential Equation (ODE)

$$v(x_t, t) = \frac{dx}{dt}(t) \approx \frac{x_{t+\Delta t} - x_t}{\Delta t}$$

FLOW MATCHING

CONTINUOUS FLOWS

FLOW MATCHING



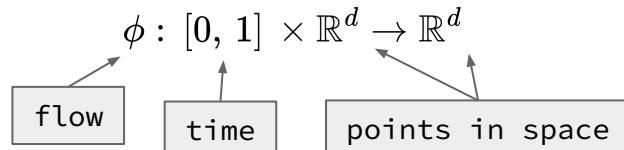
explained

Note that we no longer have to specify:

- the structure for in-between deformations
- the number of ‘integration steps’ N

This **extra freedom** is why people started considering flows

So the idea now is to consider **continuous time** deformations:



And learn vector fields that govern these deformations:

$$\frac{d\phi(t, x(t))}{dt} = u(t, x(t))$$

Simplified notation

$$\iff \frac{d\phi_t(x_t)}{dt} = u_t(x_t)$$

PROBABILITY FLOW: EXAMPLE

$X \sim \mathcal{N}(0, 1)$, $Y = \Phi(X) = X + s \sin(tX)$, where $s, t \in \mathbb{R}$ and $|st| < 1$

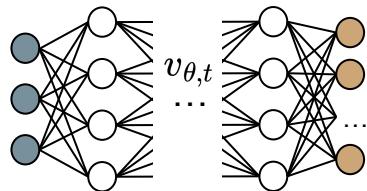
[Interactive app link](#)



Fix s and think
of t as time

CONTINUOUS NORMALIZING FLOWS (CNF)

Goal: learn the vector field which defines the flow



$$\frac{d\phi_t(x_t)}{dt} = u_t(x_t)$$

We could try to still use the following relation:

$$p_t(x) = p_{\text{init}}(\phi_t^{-1}(x)) \cdot |\det D\phi_t^{-1}(x)|$$

Probability path

to maximise the cross-entropy again:

$$\mathcal{L}(\theta) = -\mathbb{E}_{X \sim q} [\log p_1(X)]$$

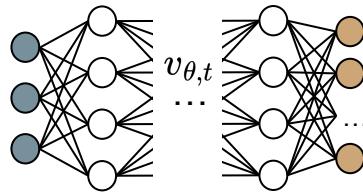
But then we have to:

- Perform numerical integration to find the final PDF of the deformation at each training step
- Somehow estimate the divergence of the vector field to compute the log (*why?*)

Long story short: **this is hard.** Instead...

CONTINUOUS NORMALIZING FLOWS (CNF)

Goal: learn the vector field which defines the flow

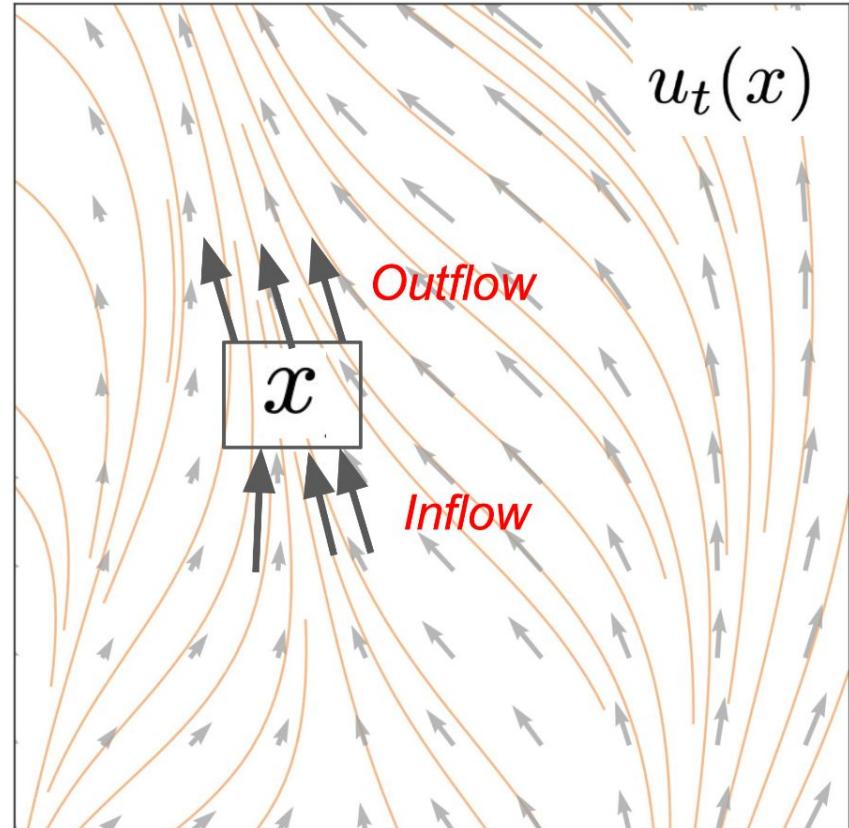


$$\frac{d\phi_t(x_t)}{dt} = u_t(x_t)$$

It turns out that if the following holds, a vector field induces a probability path (deformation):

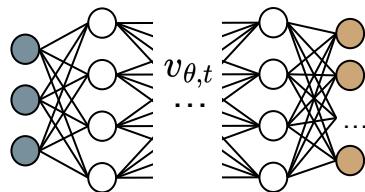
$$\frac{\partial p_t(x_t)}{\partial t} = \nabla \cdot u_t(x_t) p_t(x_t)$$

Continuity equation



FLOW MATCHING

Goal: learn the vector field which defines the flow



$$\frac{d\phi_t(x_t)}{dt} = u_t(x_t)$$

It turns out that if the following holds, a vector field induces a probability path (deformation):

$$\frac{\partial p_t(x_t)}{\partial t} = \nabla \cdot u_t(x_t) p_t(x_t)$$

Just optimise distance between the vector fields?

$$\mathcal{L}_{FM}(\theta) = \mathbb{E}_{t, p_t(x)} \|v_{\theta,t}(X) - u_t(X)\|^2$$

uniform time distribution

Vector field deforming probability correctly

Induced probability path

Neural network approximation

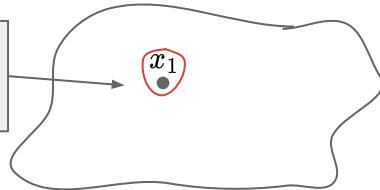
But we don't have access to neither a valid probability path nor the vector field u ...

... it's also hard to construct / propose anything reasonable

CONDITIONAL FLOW MATCHING

A trivial case when we would know how to construct a vector field is if the true distribution is *~deterministic*:

Our samples are always very similar to a single data point

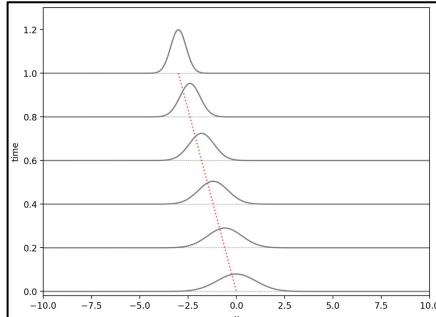


The vector field would have to simply deform init distribution into a distribution gathered around the data point, e.g. linearly:

$$p_1 = \mathcal{N}(x_1, \sigma_{\text{small}})$$



$$p_{\text{init}} = \mathcal{N}(0, 1)$$

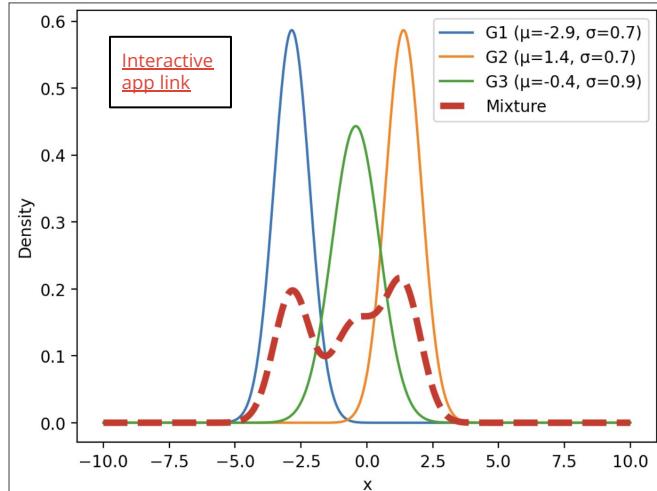


Let's rewrite a non-trivial distribution as a mixture of such simple distributions

$$p_1(x) = \int p_1(x|x_1)q(x_1)dx_1 \approx q(x)$$

We hope!

But should be true if we choose conditional probabilities well **and have enough data**



CONDITIONAL FLOW MATCHING

QUESTION: Assuming having access to all possible data points, what conditional distributions would always give us exactly the real distribution here?

ANSWER: Dirac delta distributions

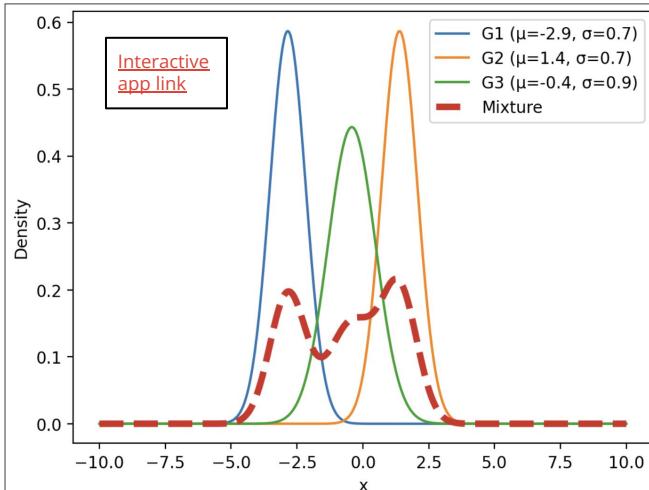
Sidenote: This procedure is kind of analogous to the discretization in space in classical physical simulations for deformable substances

Let's rewrite a non-trivial distribution as a mixture of such simple distributions

$$p_1(x) = \int p_1(x|x_1)q(x_1)dx_1 \approx q(x)$$

We hope!

But should be true if we choose conditional probabilities well **and have enough data**



CONDITIONAL FLOW MATCHING

This works at every moment on the probability path:

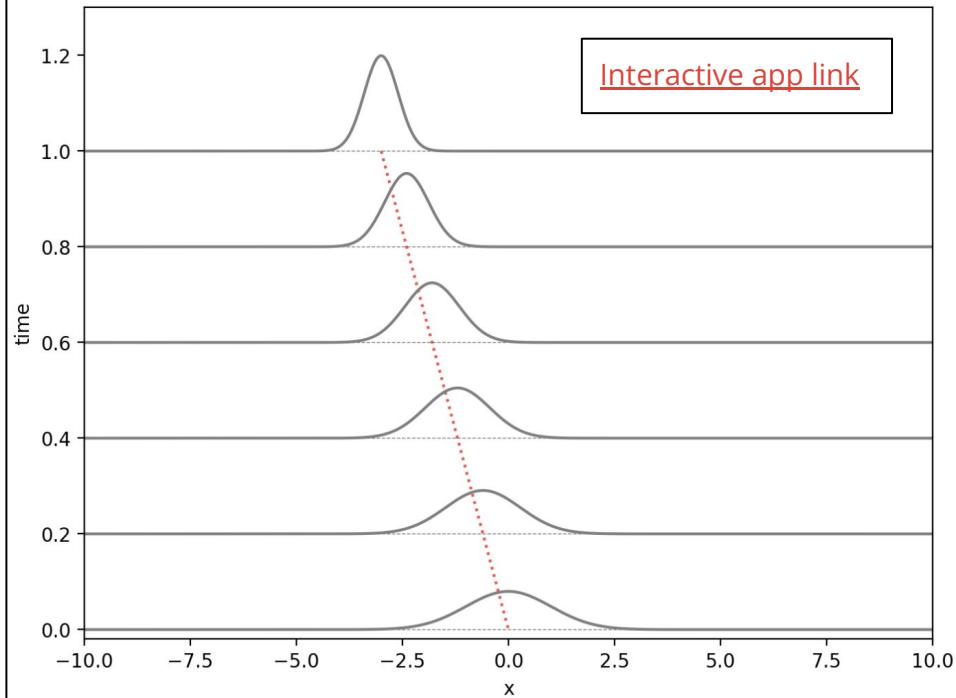
$$p_t(x) = \int p_t(x|x_1)q(x_1)dx_1$$

We end up with conditional probability paths

As before, we have a related conditional vector field that determines conditional deformation (continuity equation):

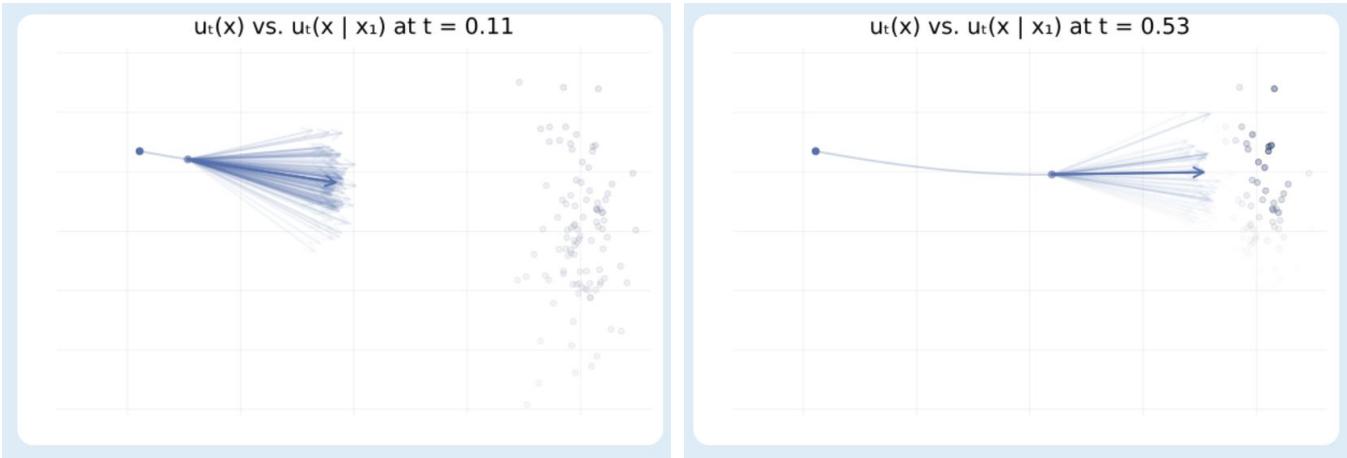
$$\frac{\partial p_t(x_t | x_1)}{\partial t} = \nabla \cdot u_t(x_t | x_1)p(x_t | x_1)$$

Example: Optimal Transport Gaussian conditional probability path.



CONDITIONAL FLOW MATCHING

What happens if we average over multiple conditional vector field directions?



In other words, we look at:

$$u_t(x_t) = \mathbb{E}_{p_{1|t}}[u_t(x_t | x_1)]$$

$$= \int u_t(x_t | x_1) \frac{p_t(x_t | x_1)q(x_1)}{p_t(x_t)} dx_1$$

$$= \int u_t(x_t | x_1) p_t(x_1 | x_t) dx_1$$

Posterior probability

One can (relatively easily) show that such a vector field induces our overall deformation:

$$\frac{\partial p_t(x_t)}{\partial t} = \nabla \cdot u_t(x_t) p_t(x_t)$$

The problem is we **almost never have closed formulas for the posterior probability (WHY?)**

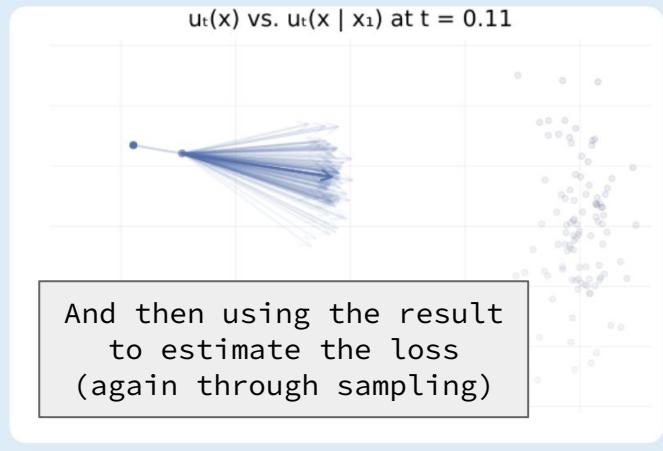
CONDITIONAL FLOW MATCHING

Surprisingly, one can (again, relatively easily) show that the unconditional loss:

$$\mathcal{L}_{CFM}(\theta) = \mathbb{E}_{t, p_t(x)} \|v_{\theta,t}(X) - u_t(X)\|^2$$

Think of changing the sampling scheme.

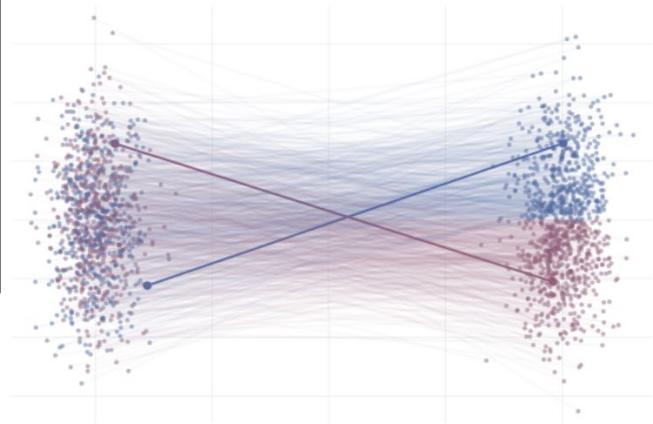
Instead
of
this:



Is up to a constant equal to a conditional loss:

$$\mathcal{L}_{CFM}(\theta) = \mathbb{E}_{t, p_t(x | x_1), q(x_1)} \|v_{\theta,t}(X) - u_t(X | X_1)\|^2$$

We sample
single
“conditional
arrows” all
over the
space



So with reasonable conditional paths and enough data, we can estimate the conditional loss and compute gradients without ever estimating the true vector field

GAUSSIAN CONDITIONAL PROBABILITY PATHS

In practice almost always gaussian conditional probability paths are used:

$$p_t(x|x_1) = \mathcal{N}(x|\mu_t(x_1), \sigma_t^2(x_1)I)$$

One can show that they have the corresponding conditional vector field

$$u_t(x|x_1) = \frac{\frac{d}{dt}\sigma_t(x_1)}{\sigma_t(x_1)}(x - \mu_t(x_1)) + \frac{d}{dt}\mu_t(x_1)$$

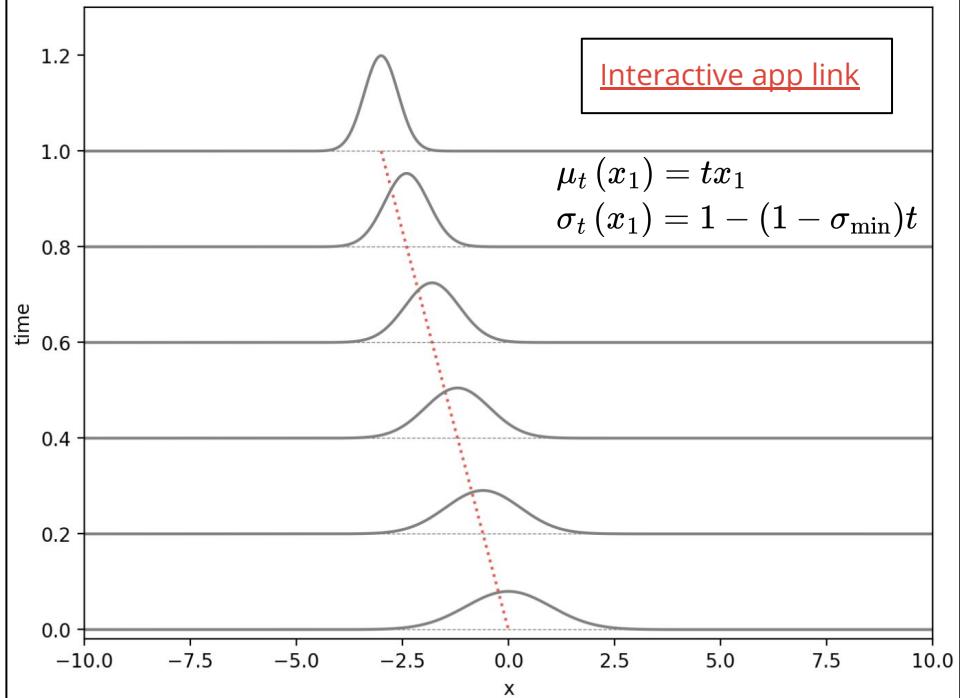
Example 2. Diffusion Process

$$\mu_t(x_1) = x_1$$

$$\sigma_t := f_\downarrow \text{ with } \sigma_1 = 0$$

Some decreasing function

Example: Optimal Transport Gaussian conditional probability path.



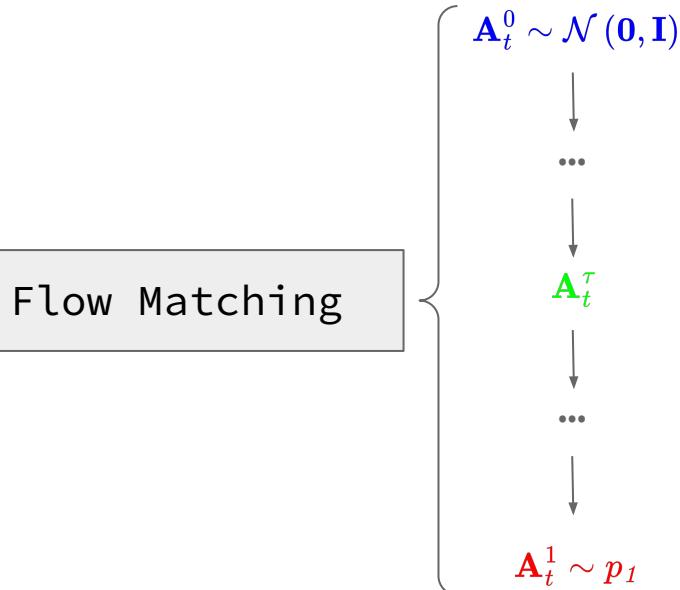
How to generate objects with a Flow Model

Algorithm 1 Sampling from a Flow Model with Euler method

Require: Neural network vector field u_t^θ , number of steps n

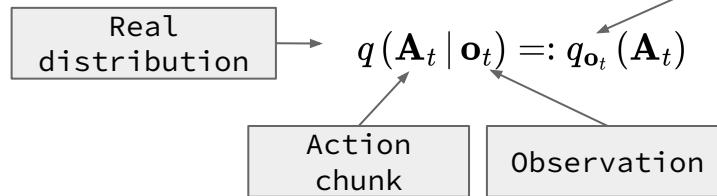
- 1: Set $t = 0$
 - 2: Set step size $h = \frac{1}{n}$
 - 3: Draw a sample $X_0 \sim p_{\text{init}}$ *Random initialization!*
 - 4: **for** $i = 1, \dots, n - 1$ **do**
 - 5: $X_{t+h} = X_t + h u_t^\theta(X_t)$
 - 6: Update $t \leftarrow t + h$
 - 7: **end for**
 - 8: **return** X_1 *Return final point*
-

EXAMPLE: PI0



FLOW MATCHING IN P10 CONTEXT

Extra notation just to easily differentiate from conditional probability paths from Flow Matching



Superscripts correspond to the flow matching timesteps, i.e. the one used for integration

$$\mathcal{L}_{CFM}(\theta) = \mathbb{E}_{t, q(x_1), p_t(x|x_1)} \|v_{\theta,t}(X) - u_t(X|X_1)\|^2$$

Subscripts denote robotic timesteps

$$\mathcal{L}^\tau(\theta) = \mathbb{E}_{q_{o_t}(\mathbf{A}_t), p(\mathbf{A}_t^\tau | \mathbf{A}_t)} \|v_\theta(\mathbf{A}_t^\tau, \mathbf{o}_t) - u_t(\mathbf{A}_t^\tau | \mathbf{A}_t)\|^2$$

Optimal Transport conditional probability paths:

$$p(\mathbf{A}_t^\tau | \mathbf{A}_t) = \mathcal{N}(\tau \mathbf{A}_t, (1-\tau)\mathbf{I})$$

Forward Euler integration to find continuous actions (starting from noise)

$$\mathbf{A}_t^{\tau+\delta} = \mathbf{A}_t^\tau + \delta \mathbf{v}_\theta(\mathbf{A}_t^\tau, \mathbf{o}_t)$$



autonomous, 1x speed

π

ARCHITECTURE OVERVIEW

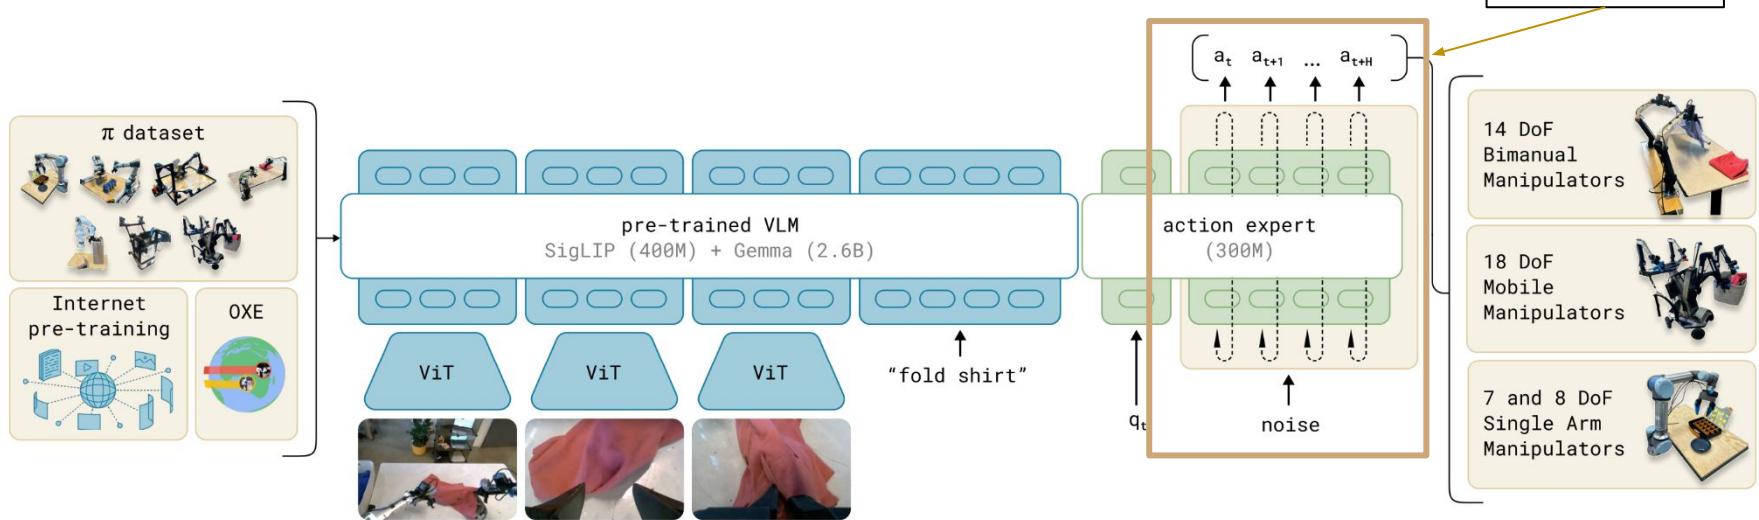
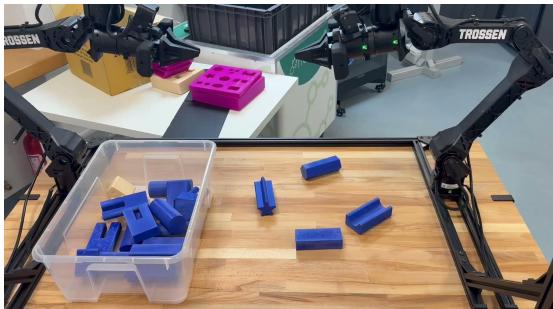
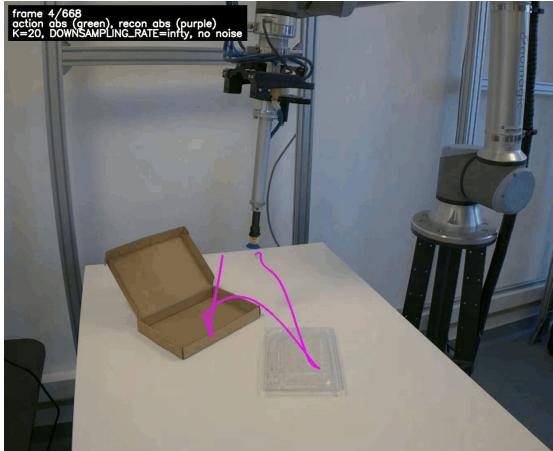


Fig. 3: **Overview of our framework.** We start with a pre-training mixture, which consists of both our own dexterous manipulation datasets and open-source data. We use this mixture to train our flow matching VLA model, which consists of a larger VLM backbone and a smaller *action expert* for processing robot states and actions. The VLM backbone weights are initialized from PaliGemma [5], providing representations learned from large-scale Internet pre-training. The resulting π_0 model can be used to control multiple robot embodiments with differing action spaces to accomplish a wide variety of tasks.

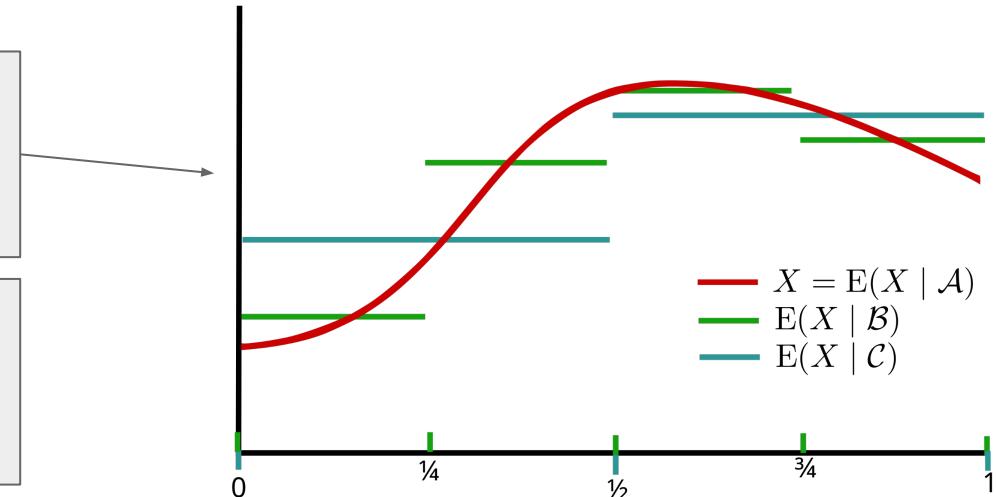
WE ARE DOING RESEARCH IN THESE AREAS AT MIMUW!



NOTE ABOUT THE PROOFS: CONDITIONAL EXPECTATION

Intuitive definition:
a **random variable** that takes values
that are the “averages” on sets which
you specify as “measurable”

A good metaphor is a camera: the higher
“resolution” sigma field (more
measurable sets) allows you to draw
more “specific” samples



With this intuition one can clearly see that:

$$\mathbb{E}[\mathbb{E}[X | \mathcal{B}] | \mathcal{A}] = \mathbb{E}[X | \mathcal{A}] \text{ whenever } \mathcal{A} \subseteq \mathcal{B}$$

We are taking a low resolution picture of a
higher resolution picture

In particular for trivial outer sigma field:

$$\mathcal{A} = \{\emptyset, \Omega\} \implies \mathbb{E}[\mathbb{E}[X | \mathcal{B}] | \mathcal{A}] = \mathbb{E}[X]$$

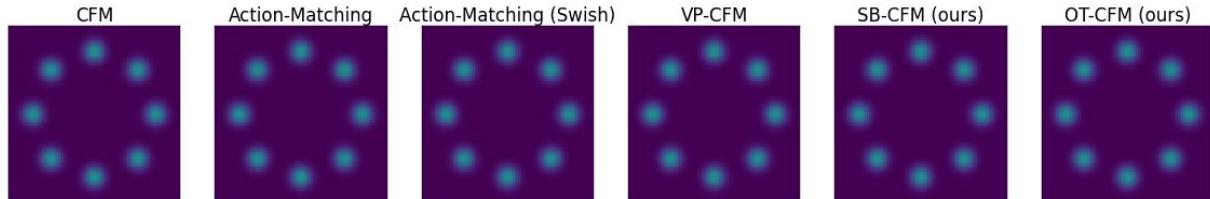
DIFFUSION

FLOW MATCHING: SUMMARY

Goal: Convert 8 two-dimensional Gaussians into “two Moons” shaped distribution

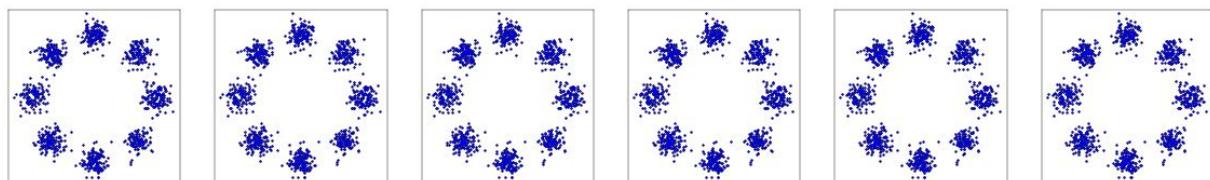
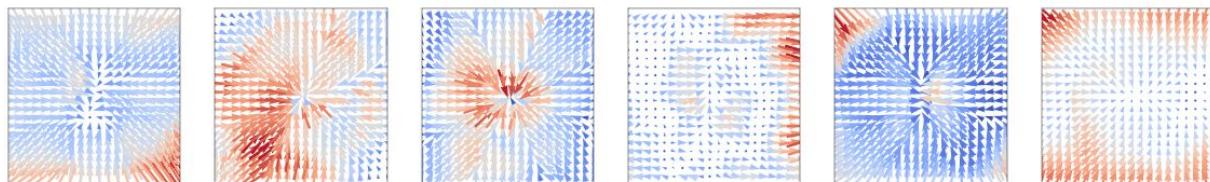
$$\phi : [0, 1] \times \mathbb{R}^d \rightarrow \mathbb{R}^d$$

flow time points in space



$$\frac{d\phi_t(x_t)}{dt} = u_t(x_t)$$

FM summarised: Numerically solve an ODE to reach the data distribution



$$\mathcal{L}_{FM}(\theta) = \mathbb{E}_{t, p_i(x)} \|v_{\theta, t}(X) - u_t(X)\|$$

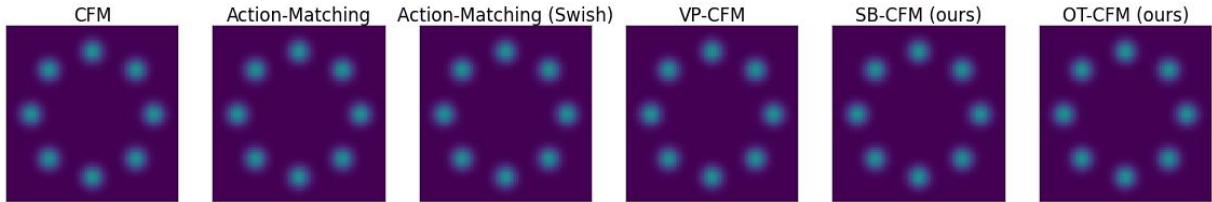
Animations source: [TorchCFM](#)

DIFFUSION: SUMMARY

Goal: Convert 8 two-dimensional Gaussians into “two Moons” shaped distribution

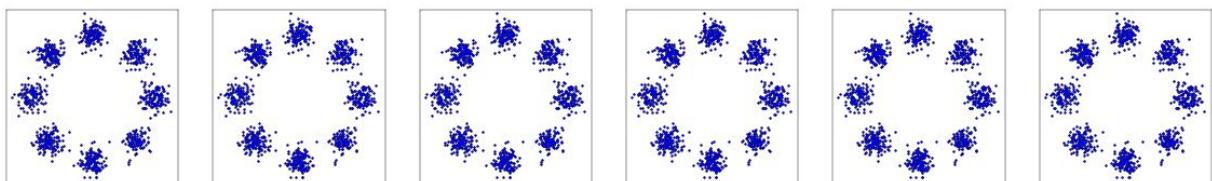
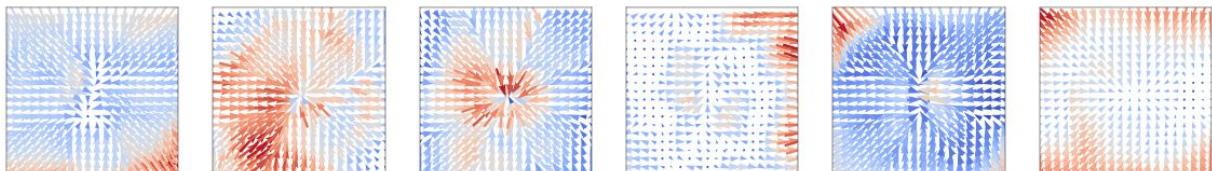
$$\phi : [0, 1] \times \mathbb{R}^d \rightarrow \mathbb{R}^d$$

flow
time
points in space



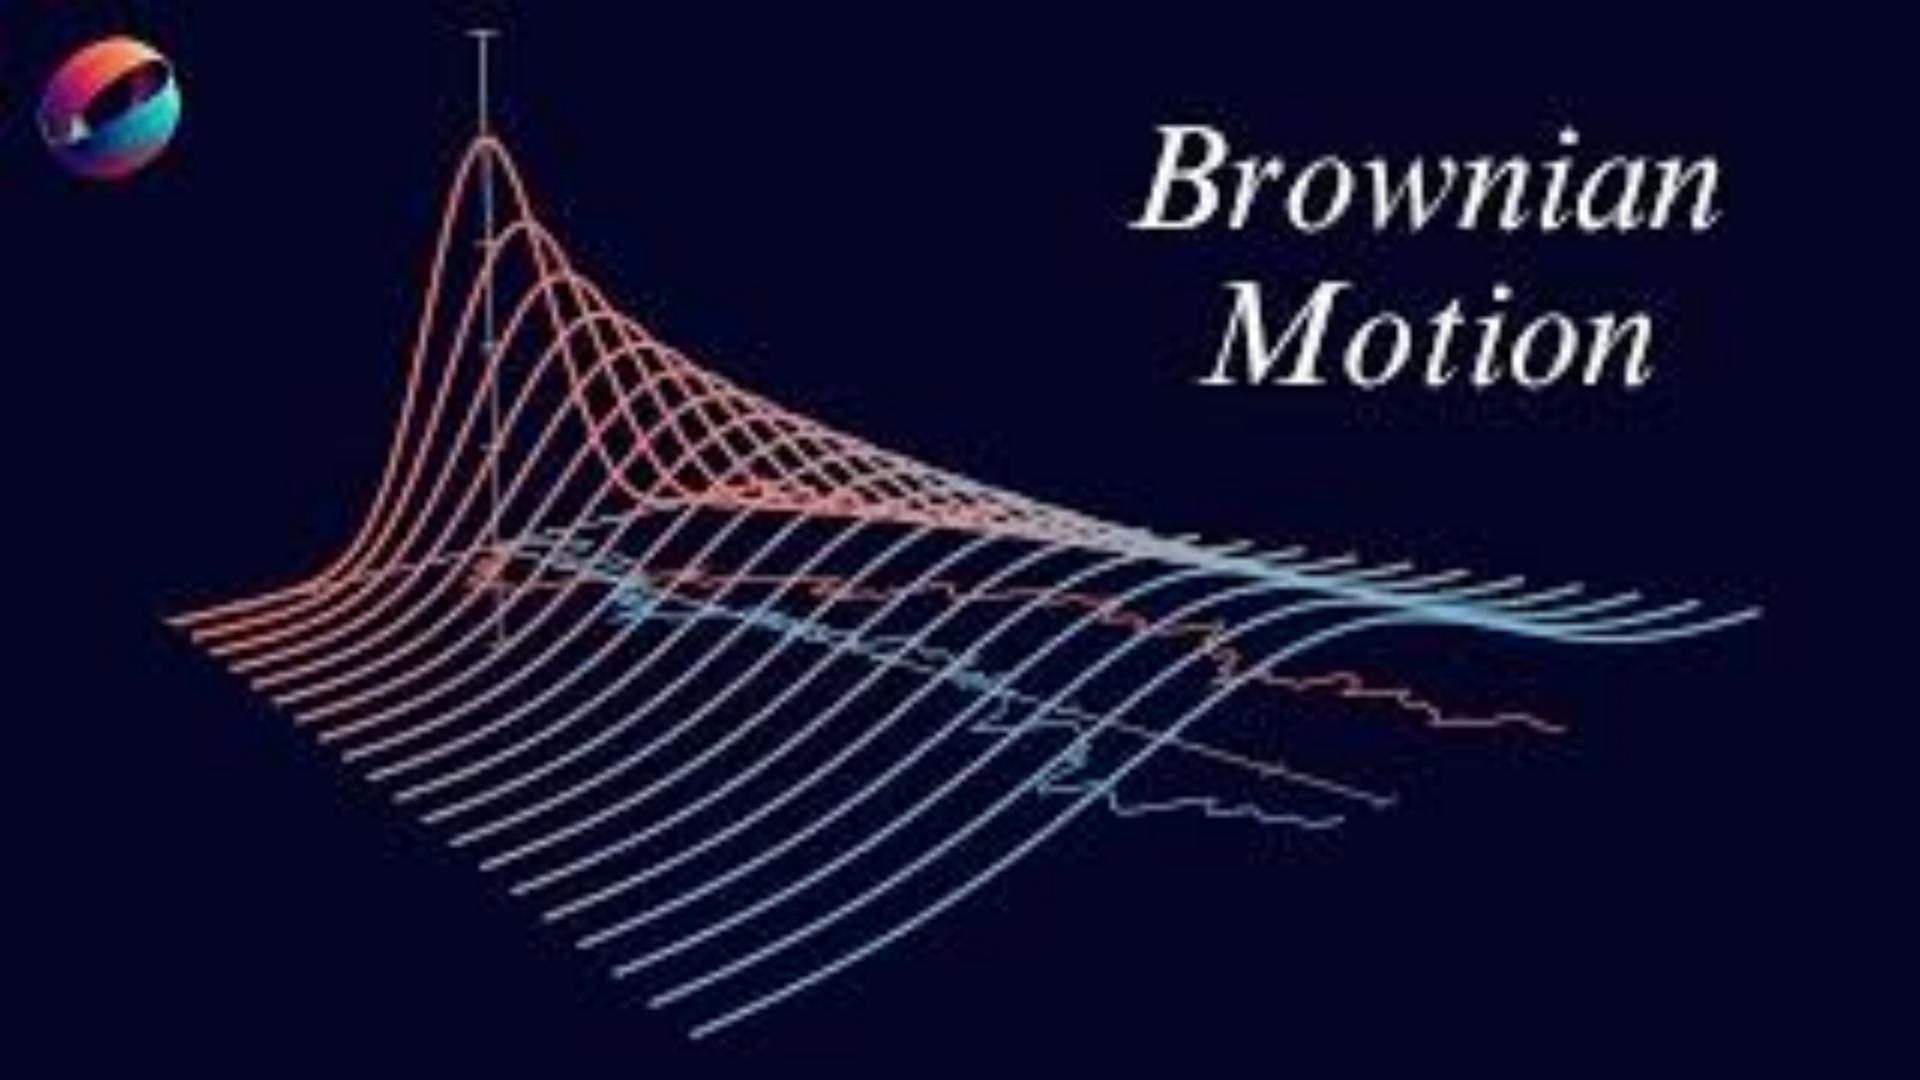
$$d\phi_t(x_t) = \frac{d\phi_t(x_t)}{dt} u_t(x_t) dt + \nabla_t(x_t) \gamma_t dW_t$$

FM summarised: Numerically solve an ODE to reach the data distribution



$$\mathcal{L}_{FM}(\theta) = \mathbb{E}_{t, p_i(x)} \|v_{\theta, t}(X) - u_t(X)\|$$

Animations source: [TorchCFM](#)



Brownian Motion



Numerical SDE simulation (Euler-Maruyama method)

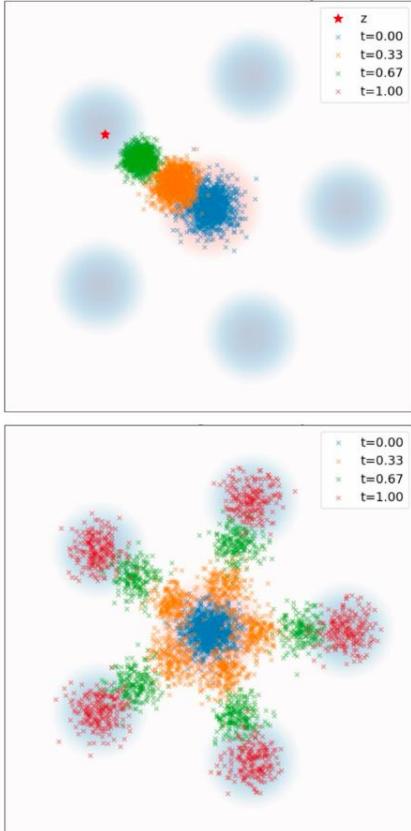
Algorithm 2 Sampling from a SDE (Euler-Maruyama method)

Require: Vector field u_t , number of steps n , diffusion coefficient σ_t

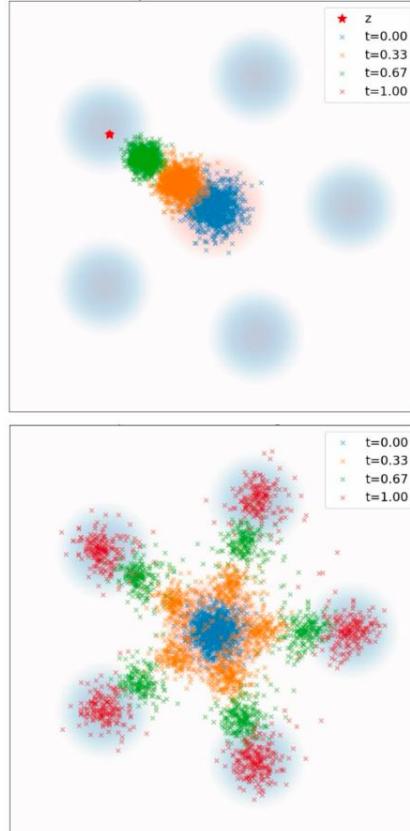
- 1: Set $t = 0$
 - 2: Set step size $h = \frac{1}{n}$
 - 3: Set $X_0 = x_0$
 - 4: **for** $i = 1, \dots, n - 1$ **do**
 - 5: Draw a sample $\epsilon \sim \mathcal{N}(0, I_d)$
 - 6: $X_{t+h} = X_t + h u_t(X_t) + \sigma_t \sqrt{h} \epsilon$ *Add additional noise with var=h scaled by diffusion coefficient σ_t*
 - 7: Update $t \leftarrow t + h$
 - 8: **end for**
 - 9: **return** $X_0, X_h, X_{2h}, X_{3h}, \dots, X_1$
-

$$p_t(\cdot | z)$$

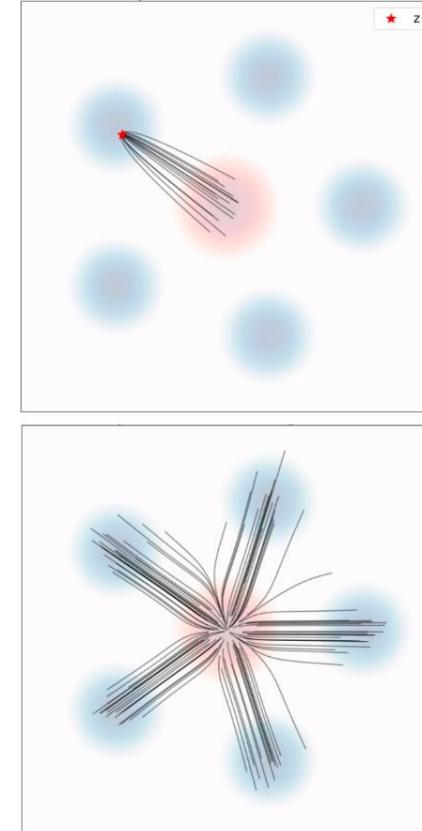
Ground truth



ODE samples

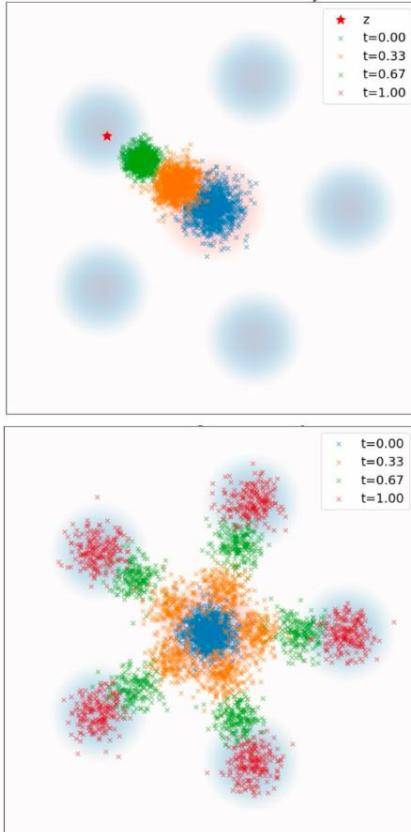


ODE Trajectories

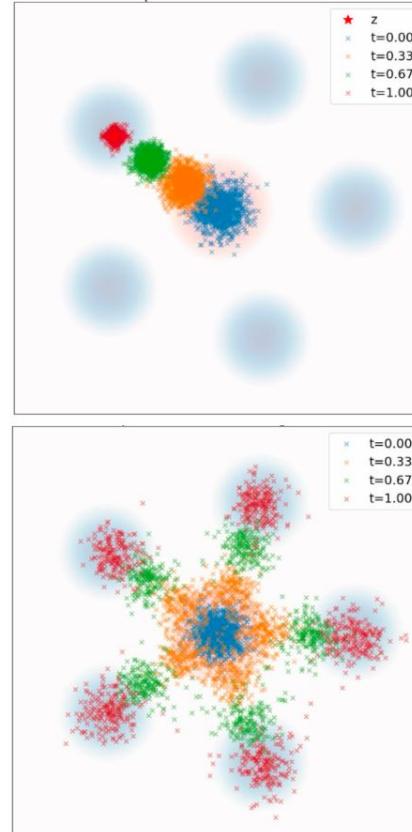


$$p_t(\cdot | z)$$

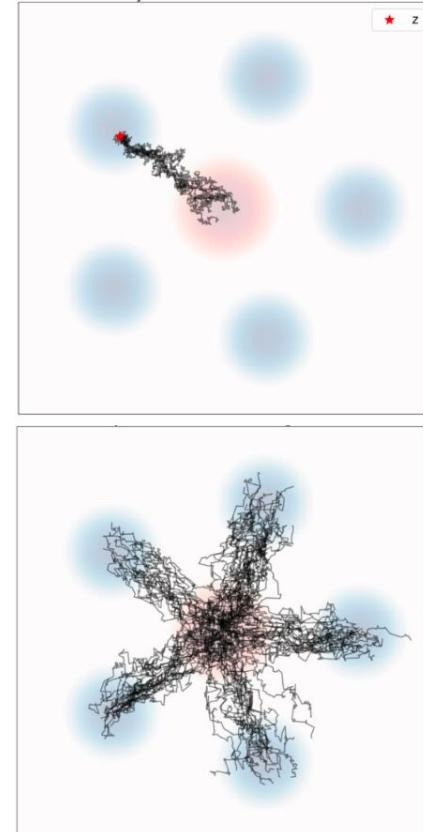
Ground truth



SDE samples

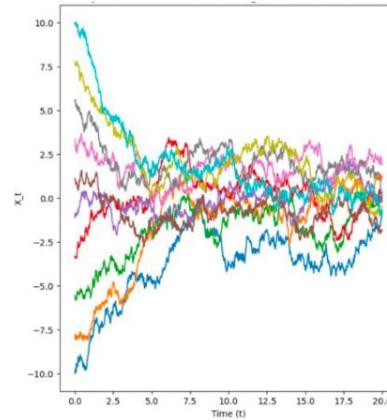
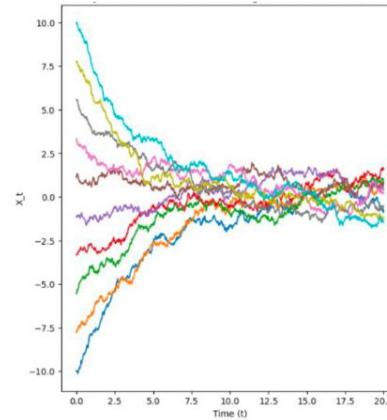
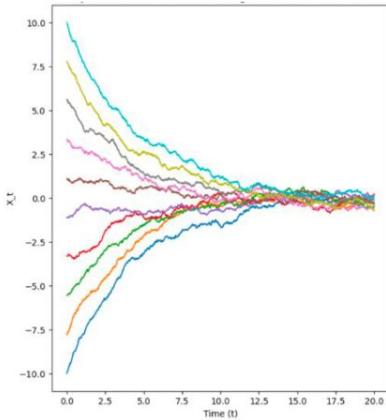
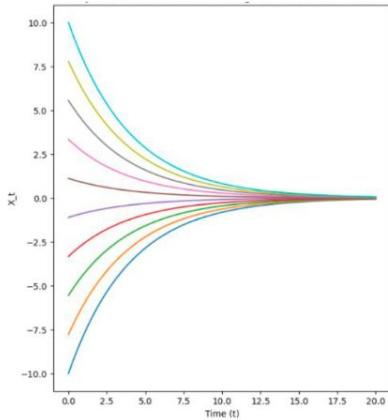


SDE Trajectories



Ornstein-Uhlenbeck Process

$$dX_t = -\theta X_t dt + \sigma dW_t$$



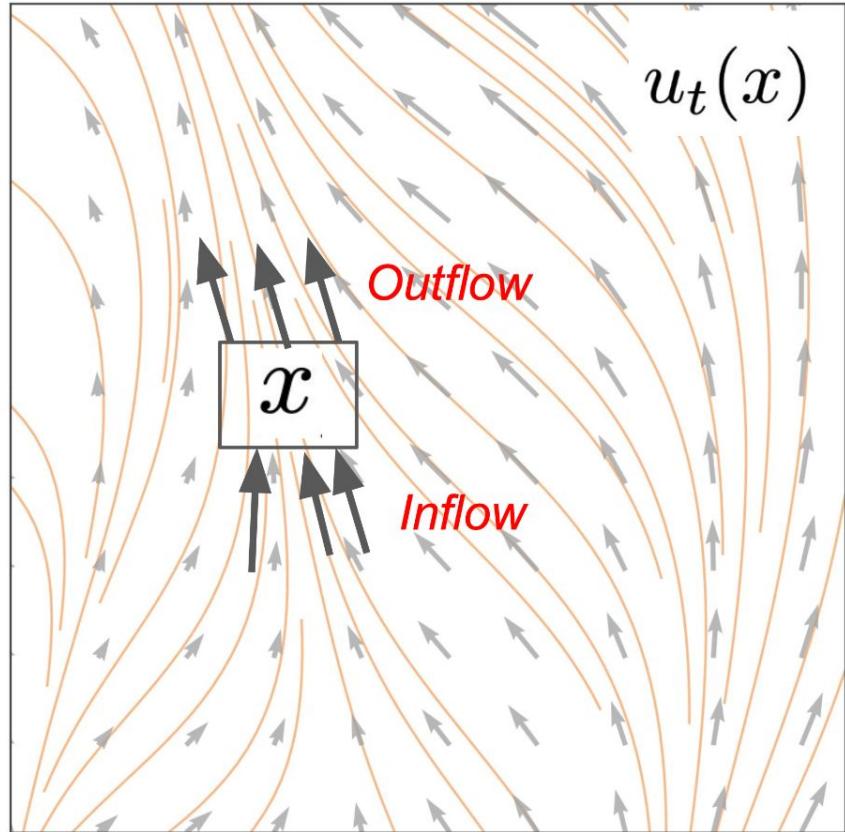
Increasing diffusion coefficient σ

Continuity Equation

$$\frac{d}{dt} p_t(x) = -\operatorname{div}(p_t u_t)(x)$$

Change of probability mass at x

Outflow - inflow of probability mass from u



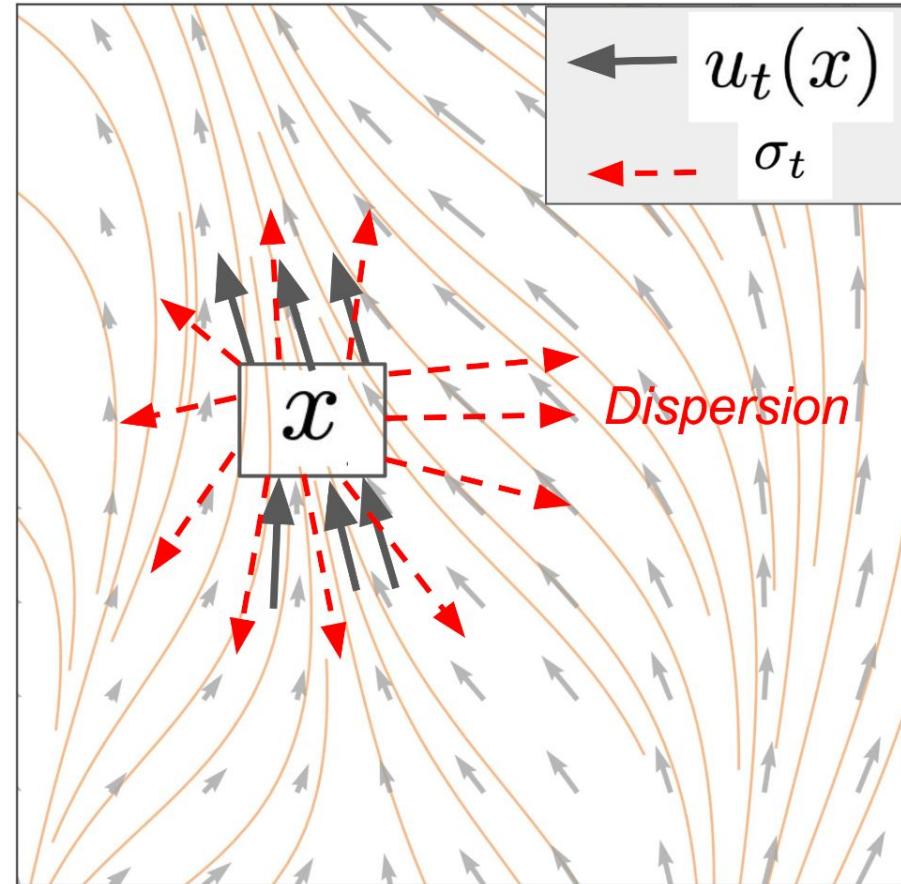
Fokker-Planck Equation

$$\frac{d}{dt} p_t(x) = -\operatorname{div}(p_t u_t)(x)$$

$$+ \frac{\sigma_t^2}{2} \Delta p_t(x)$$

Change of probability mass at x

Heat dispersion



Outlook (Next class) - Score Matching Loss

The Score Matching loss is a mean squared error between the neural network and the marginal score function:

$$L_{\text{sm}}(\theta) = \mathbb{E}_{z \sim p_{\text{data}}, x \sim p_t(\cdot|z)} [\|s_t^\theta(x) - \nabla \log p_t(x)\|^2]$$

Why this score
here?

To train a diffusion model, we need to train the score network by minimizing the score matching loss (How? Next class!)

Reminder - Marginal score function:

$$\nabla \log p_t(x) = \int \nabla \log p_t(x|z) \frac{p_t(x|z)p_{\text{data}}(z)}{p_t(x)} dz$$

*↑
Conditional score function*

SDE extension trick: The marginal score function allows to extend the ODE

to a SDE with arbitrary diffusion coefficient:

It's the correction we have to make to follow the path in presence of noise

$$X_0 \sim p_{\text{init}}, \quad dX_t = \left[u_t^{\text{target}}(X_t) + \frac{\sigma_t^2}{2} \nabla \log p_t(X_t) \right] dt + \sigma_t dW_t$$

$$\Rightarrow X_t \sim p_t$$

“Go where the log-likelihood increases the most!”

EXTRA STUFF

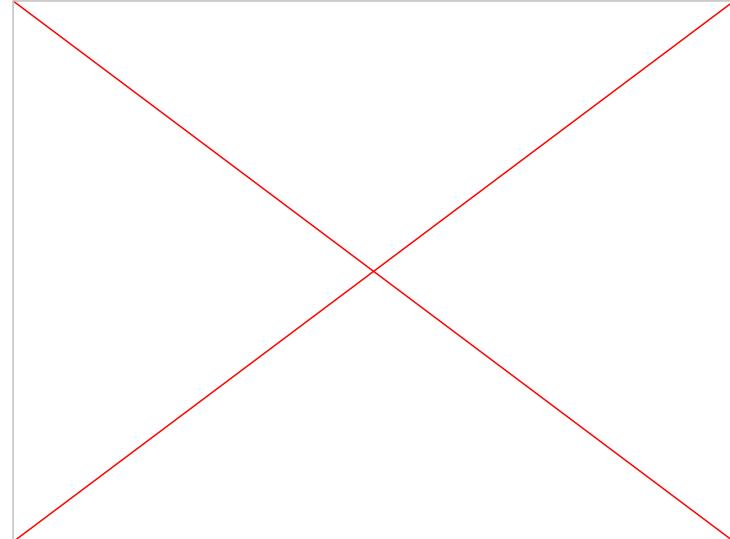
IN-PAINTING

You can run the inference in parallel to the robot execution thread.

That's why we no longer have jerky movement like this:



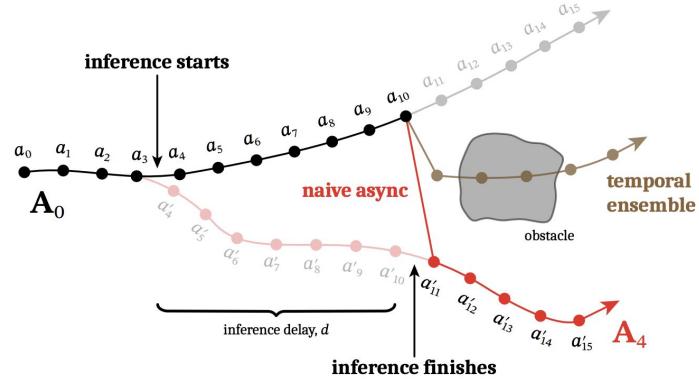
But smooth movement like this:



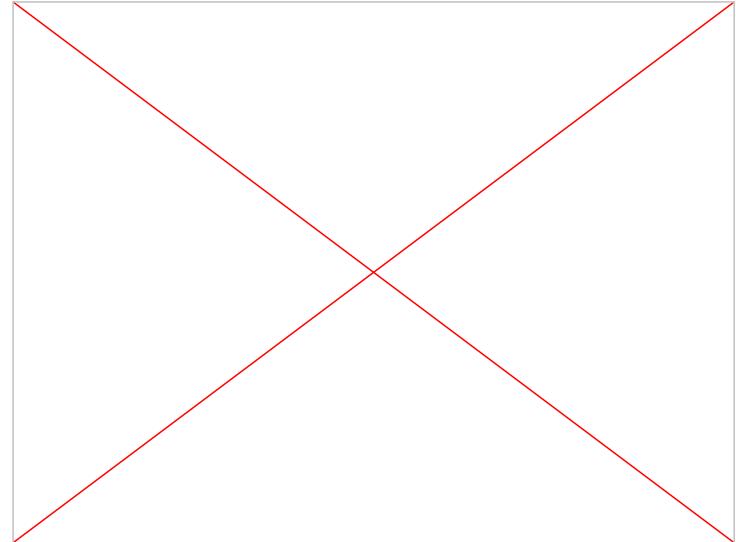
IN-PAINTING

You can run the inference in parallel to the robot execution thread.

But this requires for the inferences to be consistent:



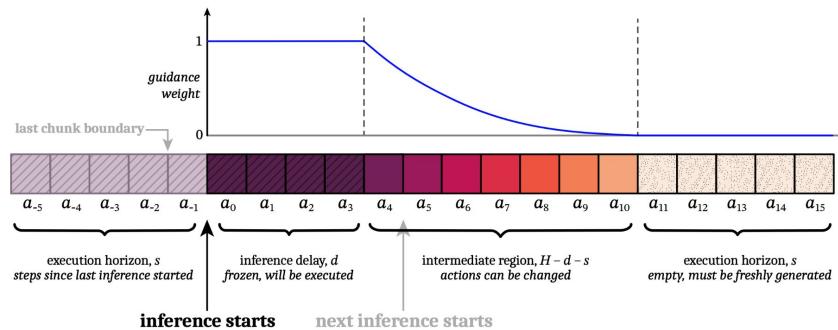
But smooth movement like this:



IN-PAINTING

You can run the inference in parallel to the robot execution thread.

But this requires for the inferences to be consistent:

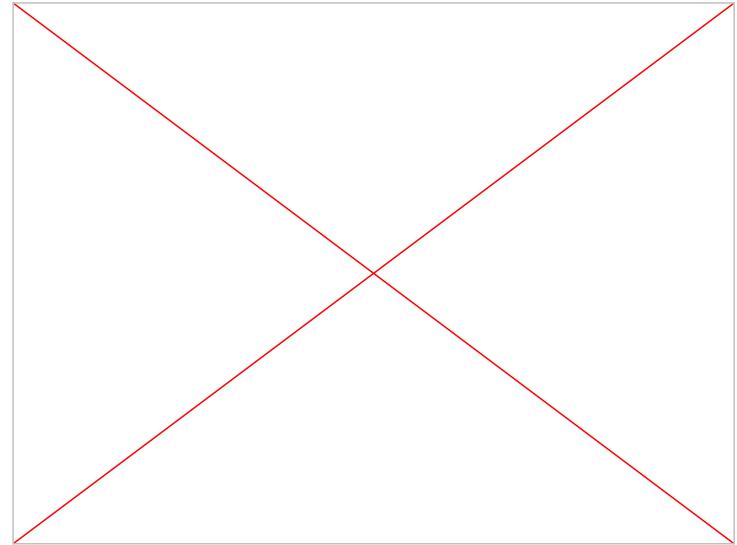


$$\mathbf{v}_{\text{IIGDM}}(\mathbf{A}_t^\tau, \mathbf{o}_t, \tau) = \mathbf{v}(\mathbf{A}_t^\tau, \mathbf{o}_t, \tau) + \min \left(\beta, \frac{1-\tau}{\tau \cdot r_\tau^2} \right) \left(\mathbf{Y} - \widehat{\mathbf{A}}_t^1 \right)^\top \text{diag}(\mathbf{W}) \frac{\partial \widehat{\mathbf{A}}_t^1}{\partial \mathbf{A}_t^\tau}$$

where $\widehat{\mathbf{A}}_t^1 = \mathbf{A}_t^\tau + (1-\tau)\mathbf{v}(\mathbf{A}_t^\tau, \mathbf{o}_t, \tau)$,

$$r_\tau^2 = \frac{(1-\tau)^2}{\tau^2 + (1-\tau)^2}.$$

But smooth movement like this:



IN-PAINTING

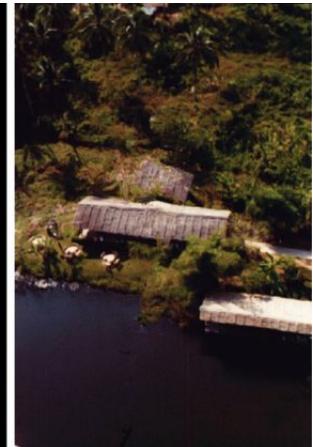
Similarly can be used for images



a



b



c

[Image source](#)

FUNCTIONAL FLOW MATCHING

Functional Flow Matching

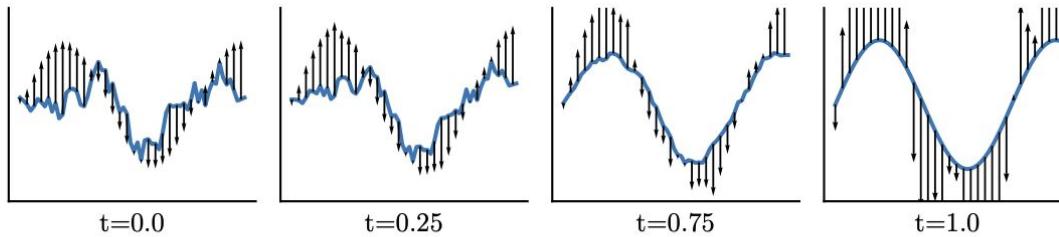


Figure 1: An illustration of our FFM method. The vector field $v_t(f) \in \mathcal{F}$ (in black) transforms a noise sample $g \sim \mu_0 = \mathcal{N}(0, C_0)$ drawn from a Gaussian process with a Matérn kernel (at $t = 0$) to the function $f(x) = \sin(x)$ (at $t = 1$) via solving a function-space ODE. By sampling many such $g \sim \mu_0$, we define a conditional path of measures μ_t^f approximately interpolating between $\mathcal{N}(0, C_0)$ and the function f , which we marginalize over samples $f \sim \nu$ from the data distribution in order to obtain a path of measures approximately interpolating between μ_0 and ν .

With proper care, FM can be generalised to work on probability measures

You can then use it to deform functions, not only finite-dimensional vectors

Hubert Lawenda's presentation during the last robot learning seminar

ADDITIONAL RESOURCES

ADDITIONAL RESOURCES

- To learn more about diffusion and flow matching models:
 - Presentation was mostly built on:
 - [An introduction to Flow Matching · Cambridge MLG Blog](#), and
 - [MIT 6.S184: Flow Matching and Diffusion Models](#) lectures
 - [TorchCFM](#) - a python package which implements Flow Matching
 - [How I Understand Flow Matching](#)
 - [Flow Matching for Generative Modeling \(Paper Explained\)](#)
 - [What are Normalizing Flows?](#)

THANK YOU