

Administrative:

- Survey - please fill survey in USOS.
- Exam free pass - threshold to be announced a few days before the exam.
- Exam - to be eligible for the exam you need to pass the homeworks:
 - 50% is guaranteed to be enough
 - Possibly lower threshold released before the exam
- Exam on Feb 8 - Saturday (5h), two parts:
 - Test (moodle, no materials allowed).
 - Programming tasks in pytorch.
 - Old exams are on on the [course website](#).

Old exams

Both coding tasks and tests from editions 2020/21-2023/24 are [here](#)

Model based RL

Plan for today:

- Exploration & Exploitation (briefly) -> AlphaGo Zero
- Model Based Reinforcement Learning (briefly) -> From GPT to ChatGPT

Exploration & exploitation

Online decision making

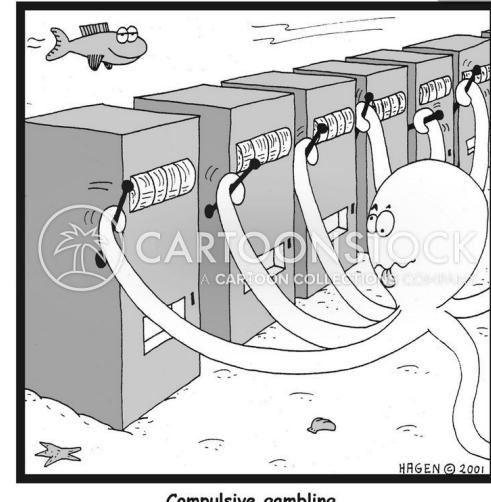
- Exploitation - make the best decision given the current knowledge
- Exploration - gather more information
- Optimum long term strategy requires a trade-off between them

Examples

- Lunch selection:
 - Exploitation - pick your favourite dish
 - Exploration - try something new
- Game playing:
 - Exploitation - play the best move (as you currently know)
 - Exploration - try an experimental move

Multi-armed bandits

- Set of bandits (actions) A
- Each turn an action is chosen independently
- The environment generates a reward $Q(a) = \mathbb{E}[r|a]$
 - **Reward depends only on most recent action!**
- Goal: maximize total rewards
 - e.g., recommender systems
 - exploration vs exploitation



Compulsive gambling

Regret minimization

- Maximizing reward is equivalent to minimizing regret
 - a^* - best action in hindsight

$$R_t = \sum_{i=1}^t (\mathbb{E}[a^*] - \mathbb{E}[r|a_t])$$

- Regret minimization is an established mathematical problem.

Greedy action selection

- Maintain Monte Carlo estimates of the action values:

$$\bar{Q}_t(a) = 1/N_t(a) \sum_{i=1}^t [r_t \text{ if } a_t = a \text{ else } 0]$$

- Greedy always picks action with highest estimate

$$a_t = \operatorname{argmax}_a \bar{Q}_t(a)$$

- Problem - greedy can use suboptimal action forever

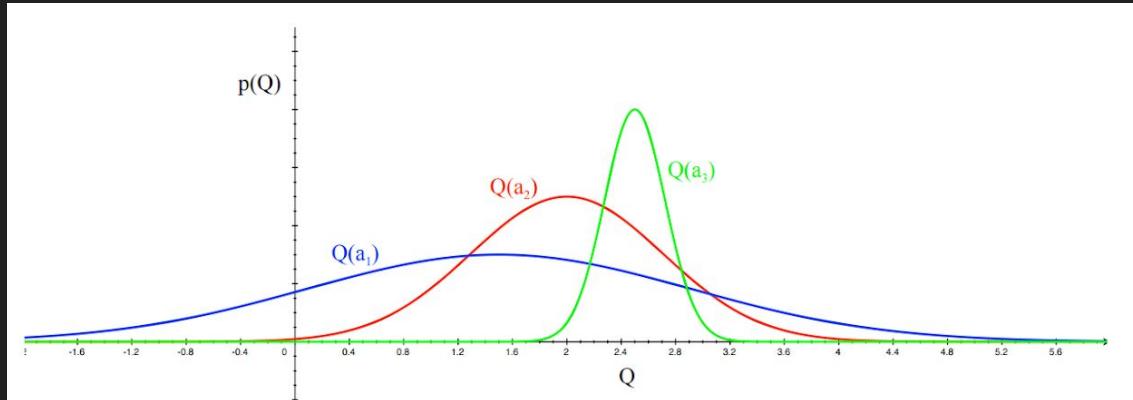
Epsilon greedy

ε -greedy exploration strategy:

- With probability ε act randomly
- With probability $(1-\varepsilon)$ act greedily
 - according to the current estimates

Modeling rewards as distributions

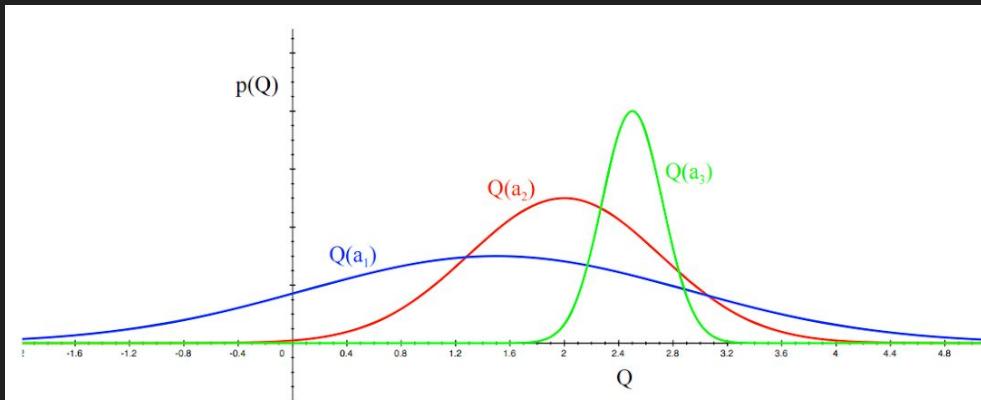
- Consider having distribution modeling current belief about the expected rewards



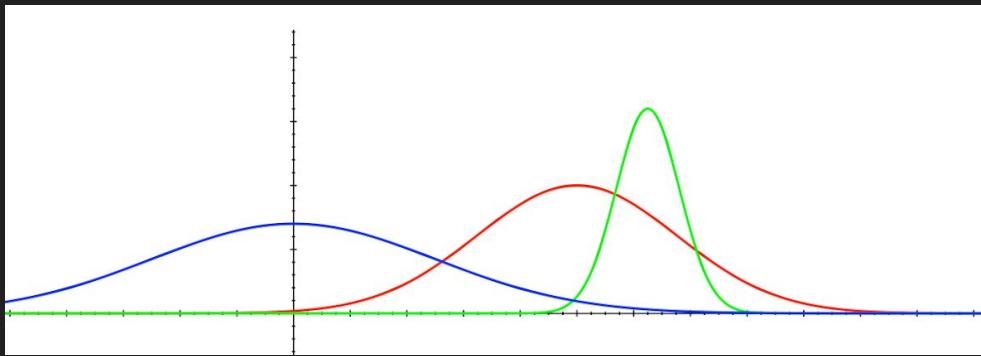
- Which action should we pick?
 - Uncertain actions are worth exploring

Modeling rewards as distributions

- Current belief:



- Updated belief
after picking blue



UCB - Upper Confidence Bounds

- Maintain uncertainty of each action $\bar{U}_t(a)$ so that with high probability $Q(a) \leq \bar{Q}_t(a) + \bar{U}_t(a)$
- Select an action maximizing UCB
$$a_t = \operatorname{argmax}_a \bar{Q}_t(a) + \bar{U}_t(a)$$
- Update bounds for **all** actions (not only the selected one)

$$\bar{U}_t(a) = \sqrt{2 \log t / N_t(a)}$$

AlphaGo Zero

Go in numbers



**3,000
Years Old**



**40M
Players**



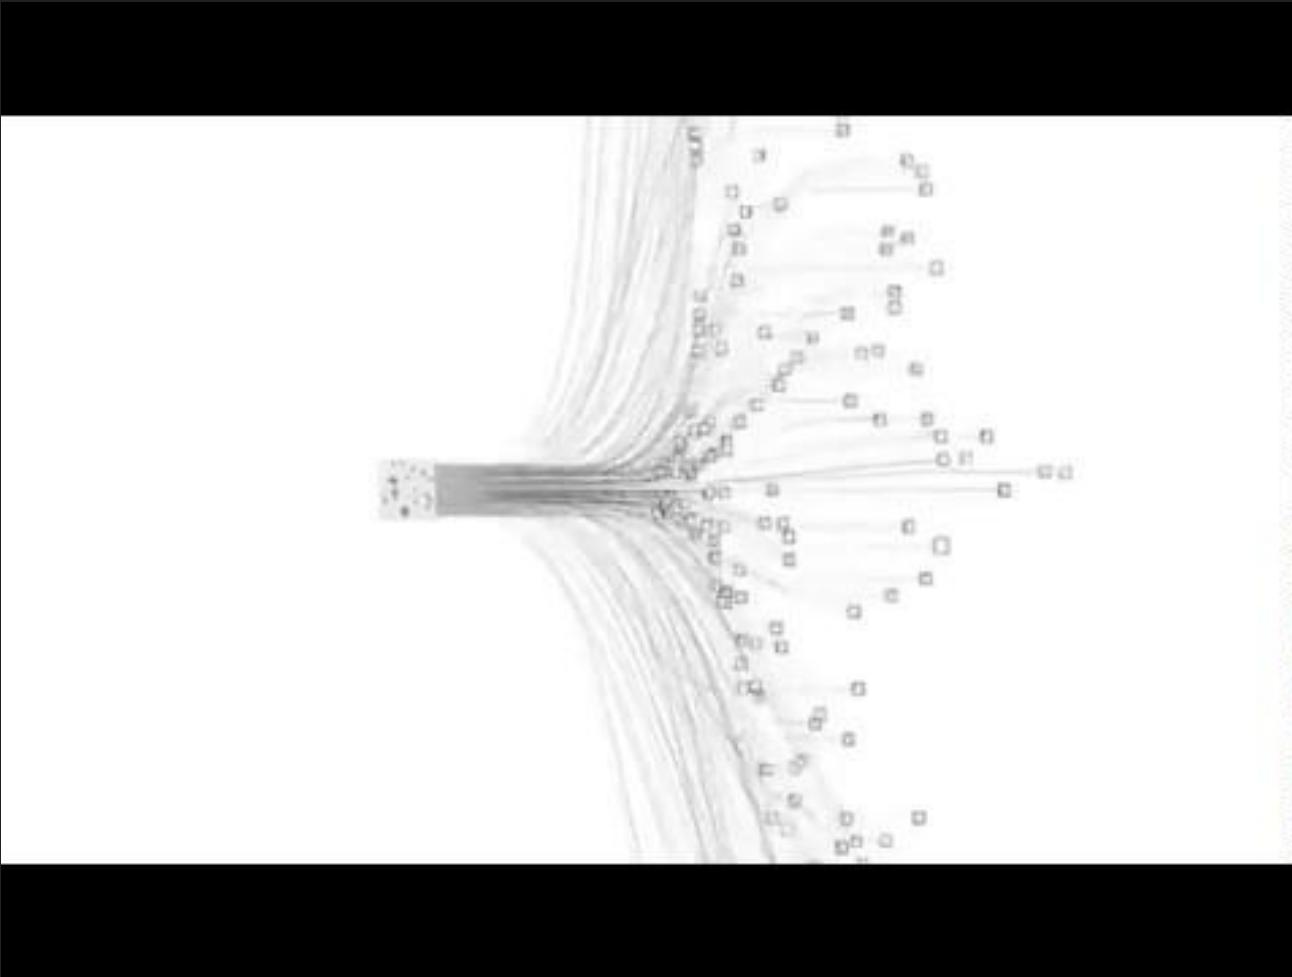
**10^{170}
Positions**



Google DeepMind



[source](#)



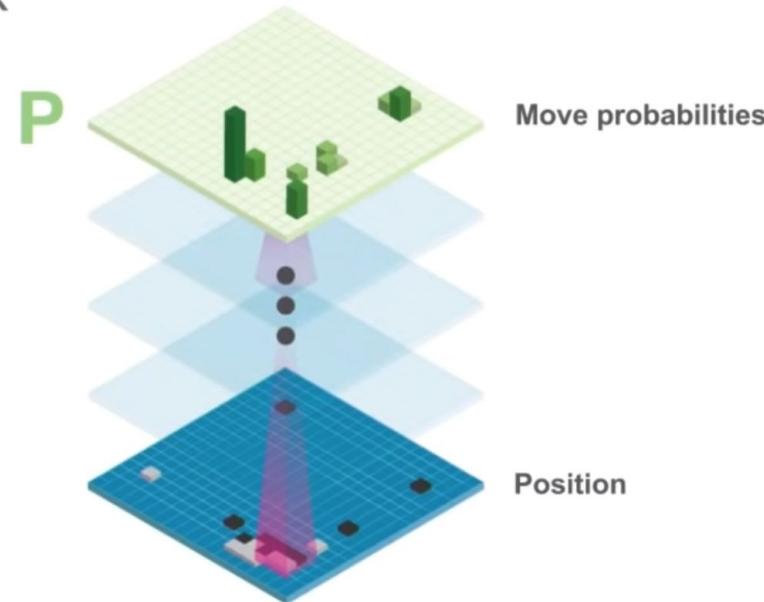
[source](#)



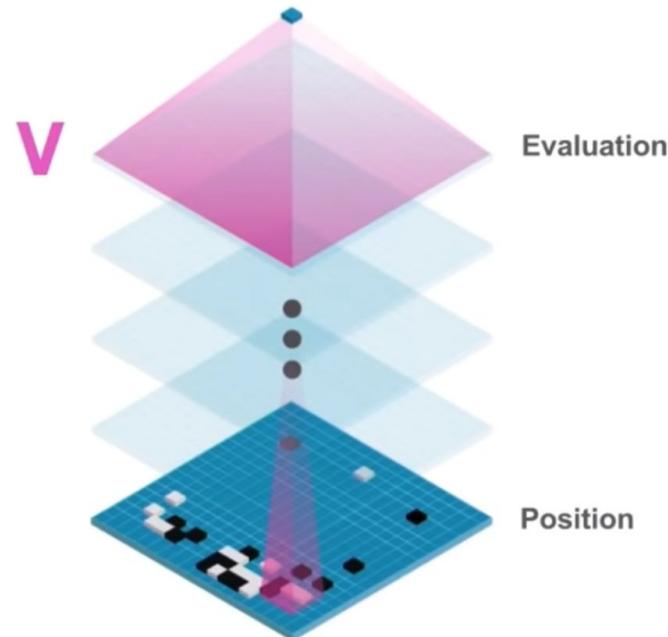
AlphaGo

First computer program to
defeat a world champion

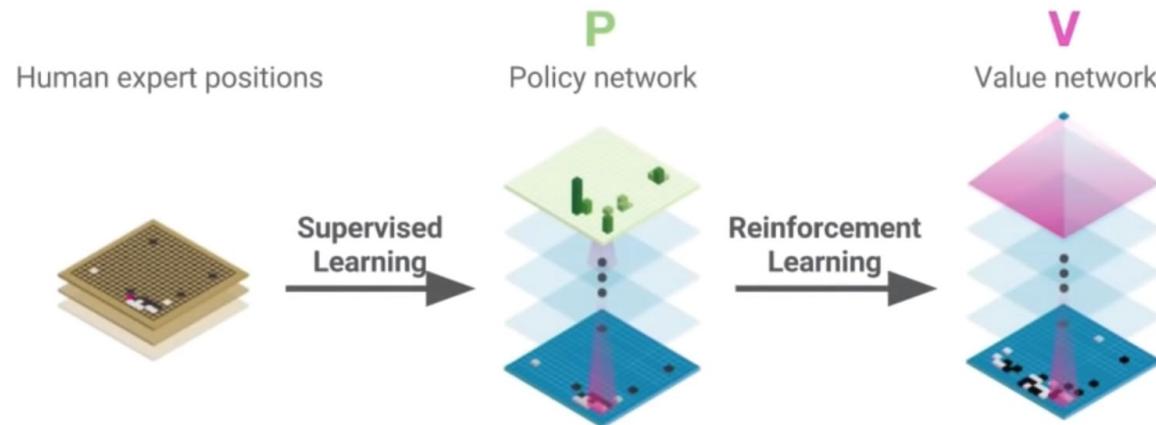
Policy network



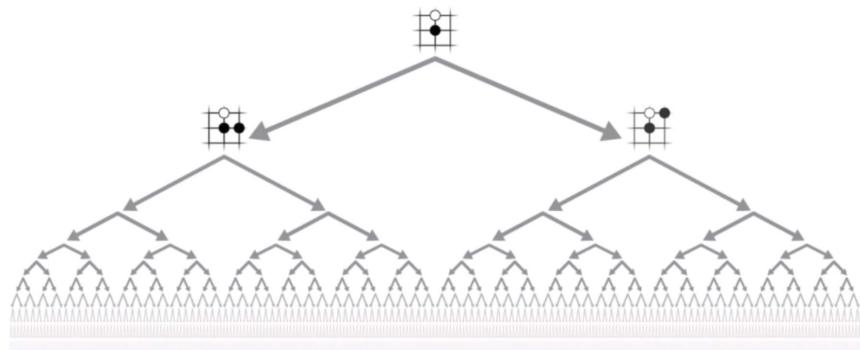
Value network



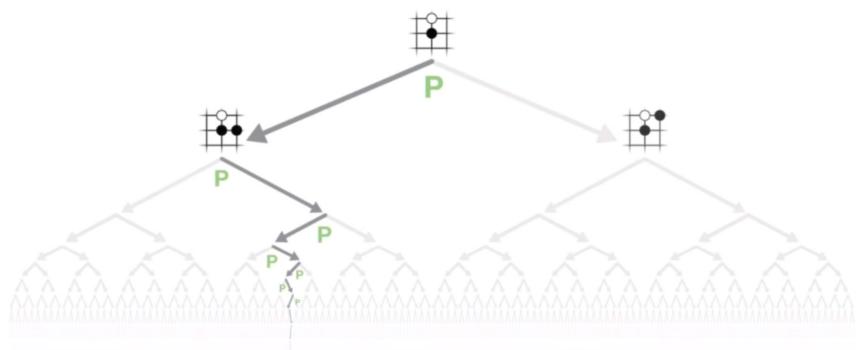
Training AlphaGo



Exhaustive search

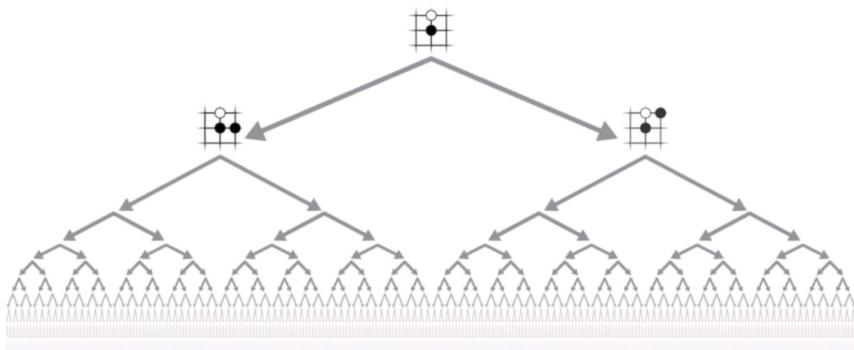


Reducing breadth with policy network

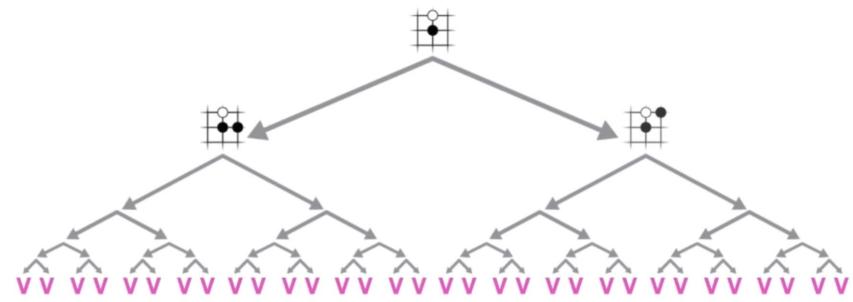


[source](#)

Exhaustive search

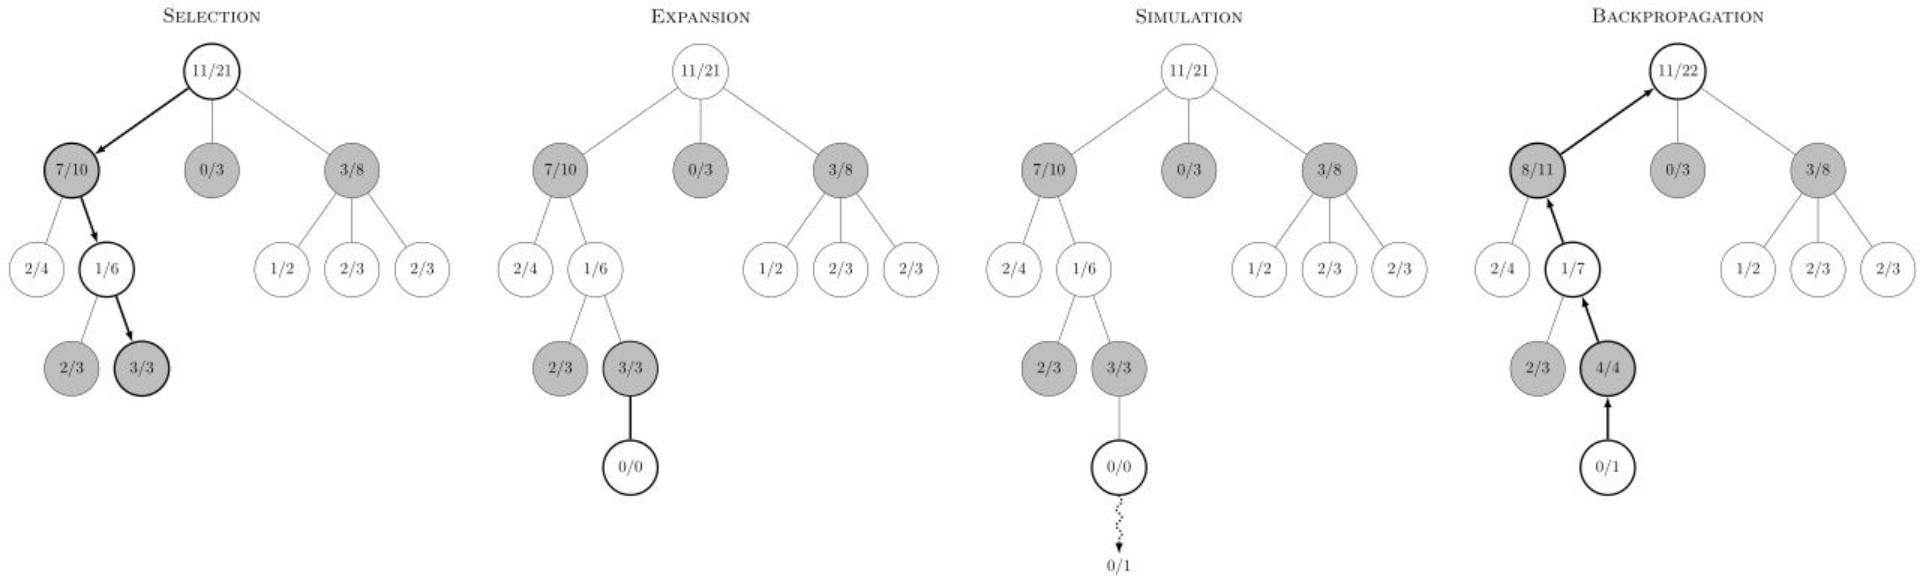


Reducing depth with value network



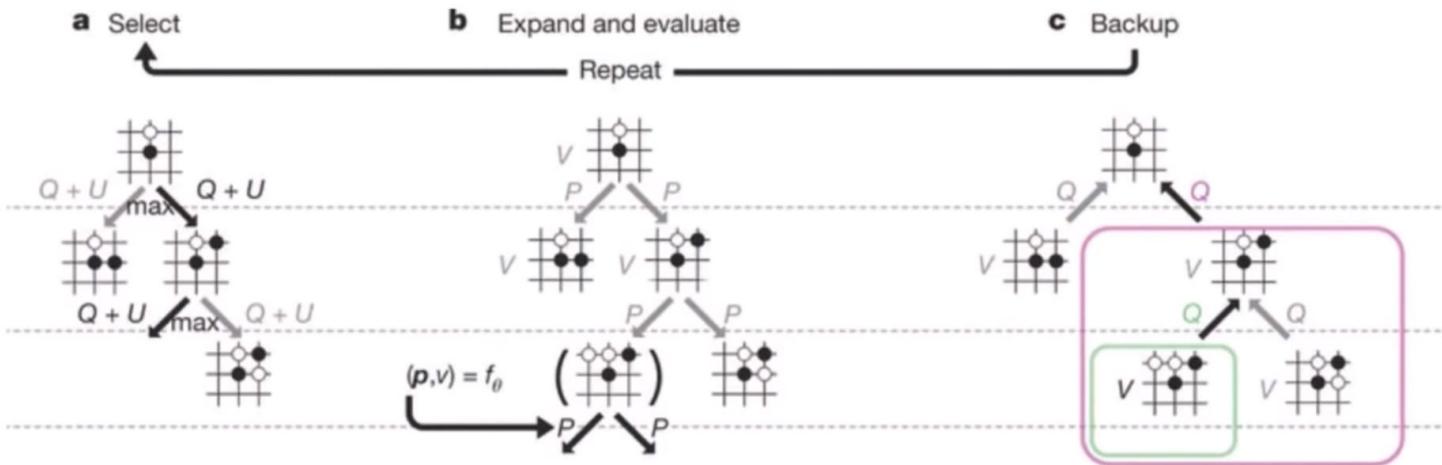
[source](#)

MCTS - Monte Carlo Tree Search



[source](#)

Monte-Carlo tree search in AlphaGo



AlphaGo vs Lee Sedol

Lee Sedol (9p): winner of 18 world titles

Match was played in Seoul, March 2016

AlphaGo won the match 4-1



AlphaGo Master vs Ke Jie

Ke Jie (9p): player ranked #1 in world

Match was played in China, May 2017

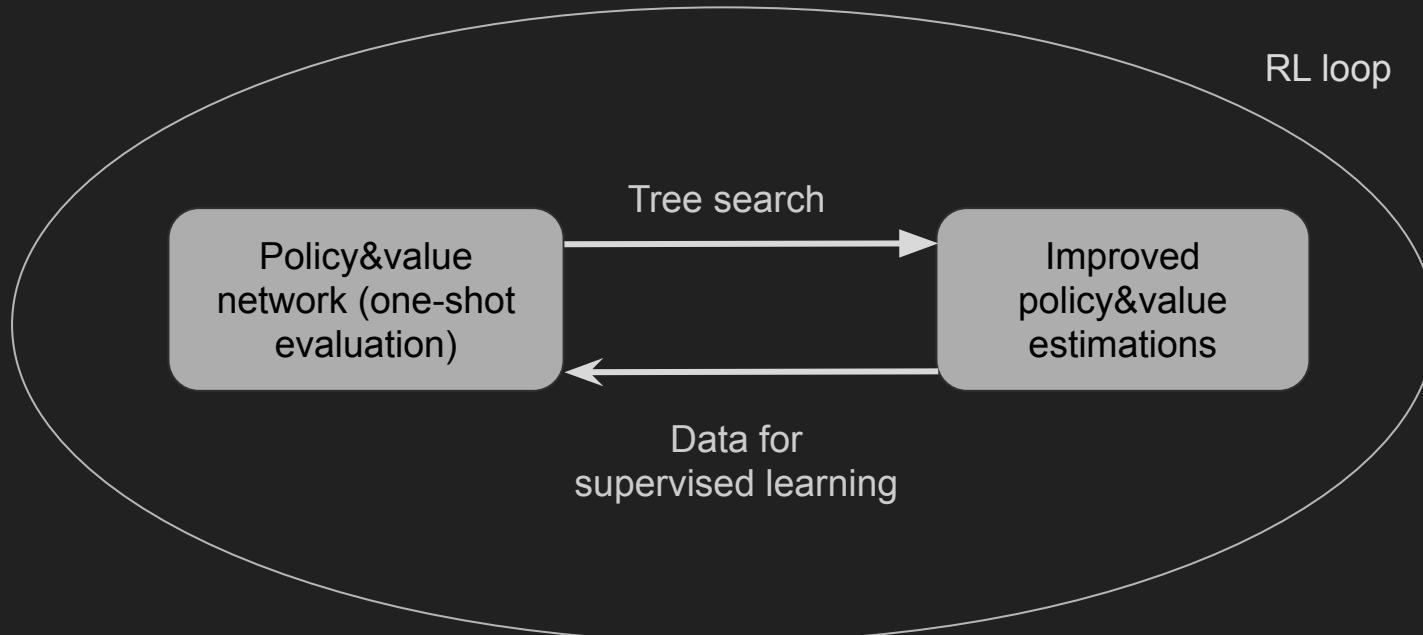
AlphaGo Master won the match 3-0



AlphaGo Zero



Mastering Go without Human Knowledge

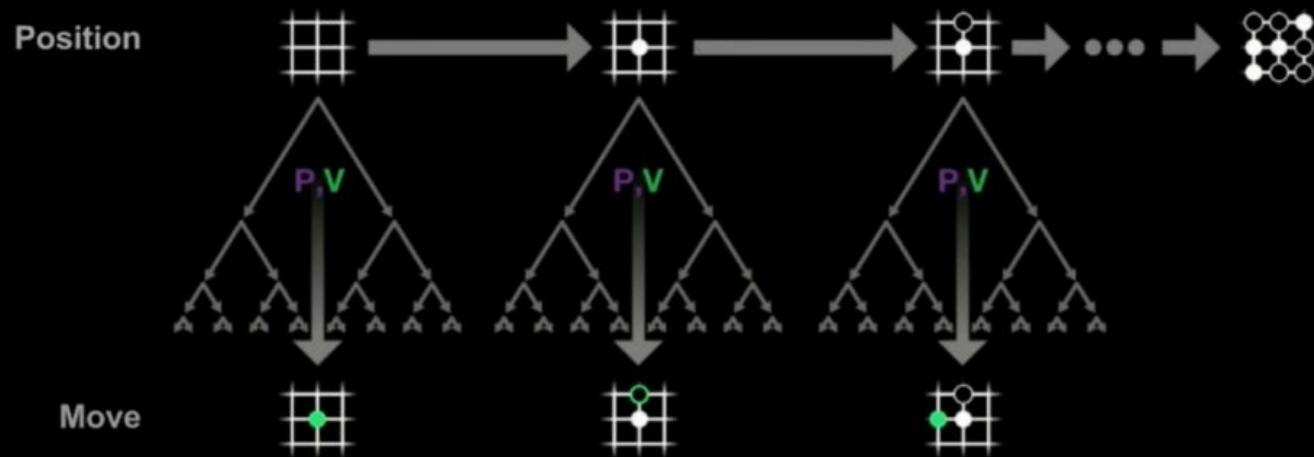


AlphaGo Zero: learning from first principles

- **No human data**
 - Learns solely by self-play reinforcement learning, starting from random
- **No human features**
 - Only takes raw board as an input
- **Single neural network**
 - Policy and value networks are combined into one neural network (resnet)
- **Simpler search**
 - No randomised Monte-Carlo rollouts, only uses neural network to evaluate

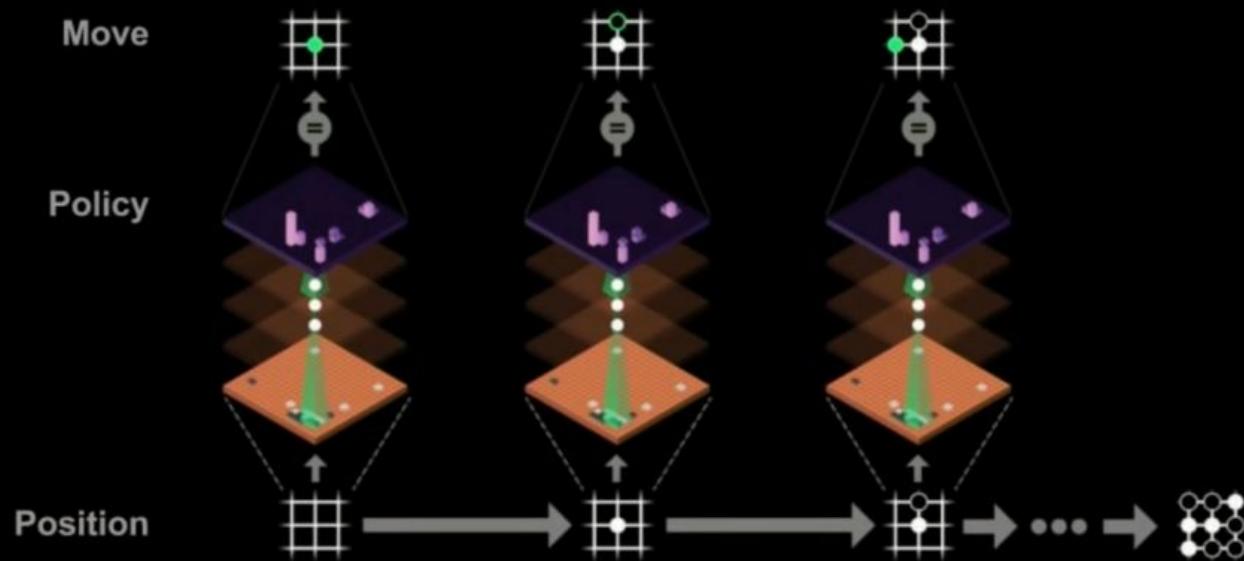
Less complexity => more generality

Reinforcement Learning in AlphaGo Zero



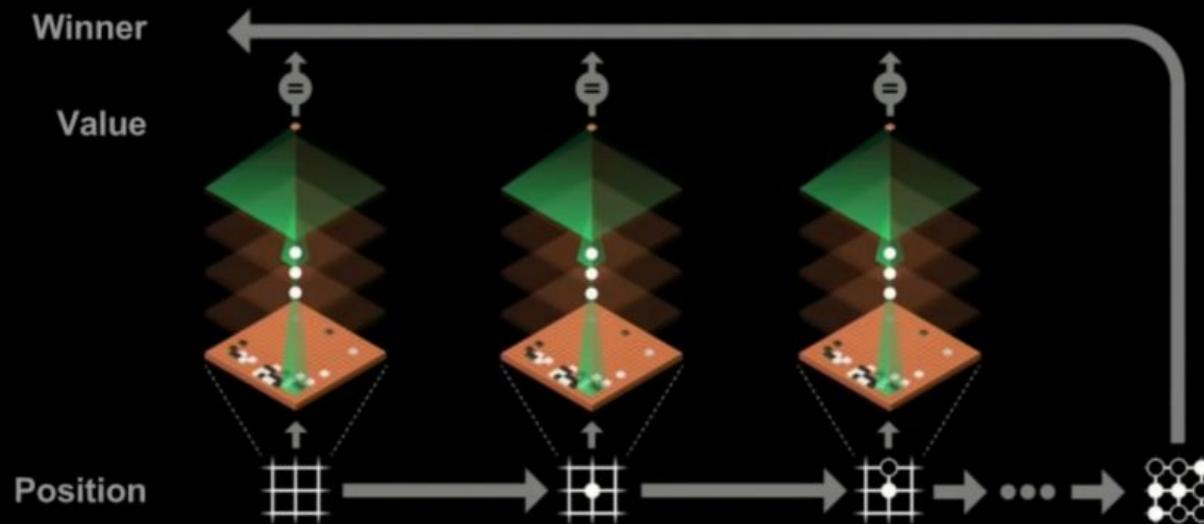
AlphaGo plays games against itself

Reinforcement Learning in AlphaGo Zero



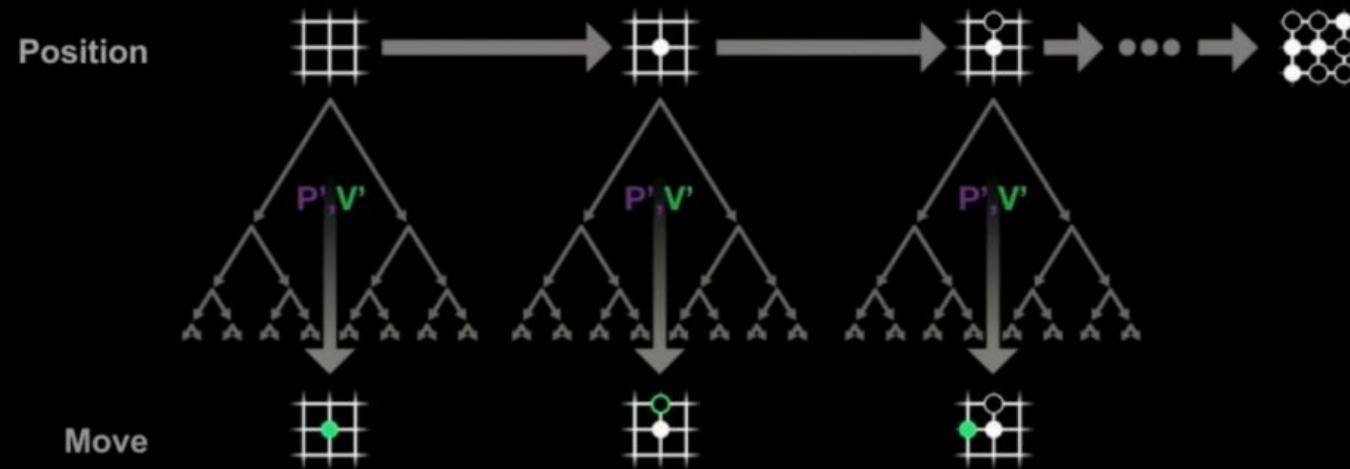
New policy network P' is trained to predict AlphaGo's moves

Reinforcement Learning in AlphaGo Zero



New value network V' is trained to predict winner

Reinforcement Learning in AlphaGo Zero



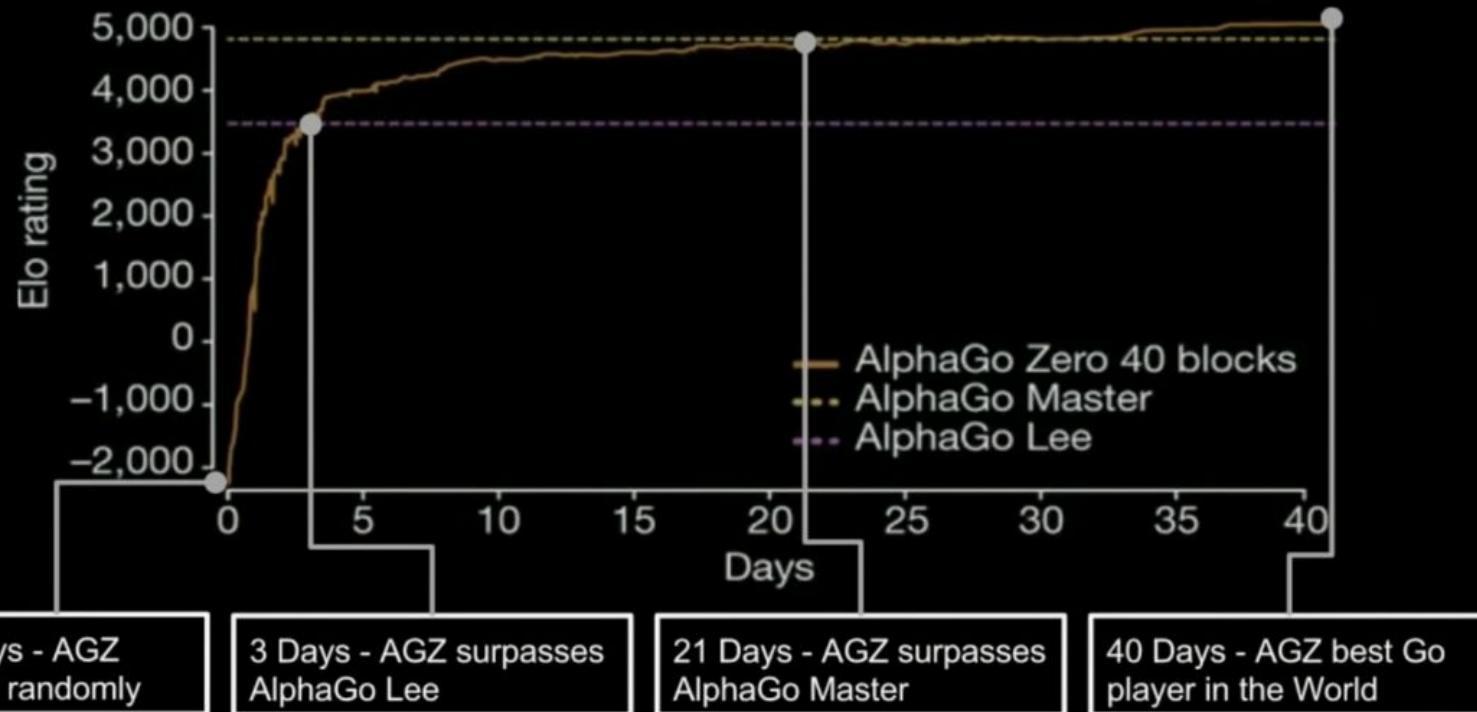
New policy/value network is used in next iteration of AlphaGo Zero

Search-Based Policy Iteration

- **Search-Based Policy Improvement**
 - Run MCTS search using current network
 - Actions selected by MCTS > actions selected by raw network
- **Search-Based Policy Evaluation**
 - Play self-play games with AlphaGo search
 - Evaluate improved policy by the average outcome

See also: Lagoudakis 03, Scherrer 15, Anthony 17

AlphaGo Zero: Learning Curve



Compute

- Alpha Go:
Policy network: 10,000 minibatches of 128 games, using 50 GPUs, 1 day.
Evaluation (MCTS): 40 search threads, 48 CPUs, and 8 GPUs
- AlphaGo Zero:
Over the course of training, 4.9 million games of self-play were generated, using 1,600 simulations for each MCTS, which corresponds to approximately 0.4s thinking time per move. Parameters were updated from 700,000 mini-batches of 2,048 positions.

*AlphaGo Zero and AlphaGo Master ... single machine with 4 TPUs;
AlphaGo Fan and AlphaGo Lee .. 176 GPUs and 48 TPUs respectively.*

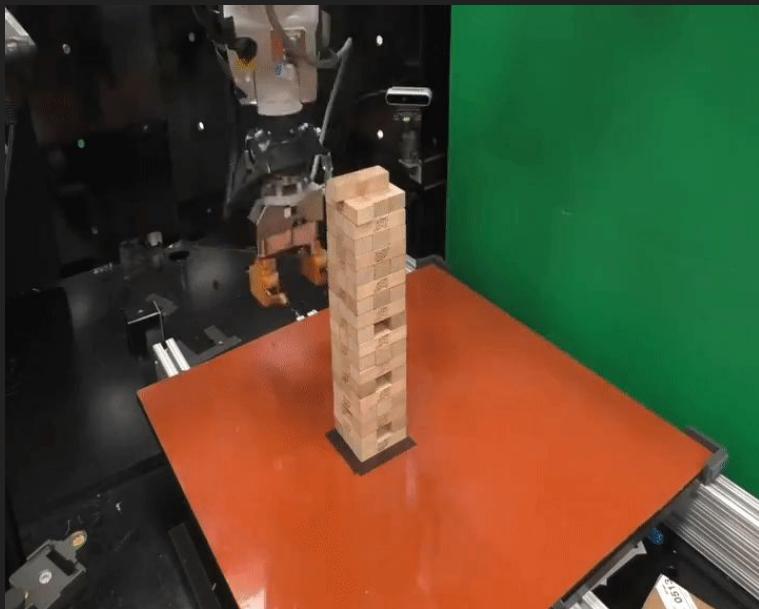
Model based RL

Learning resources

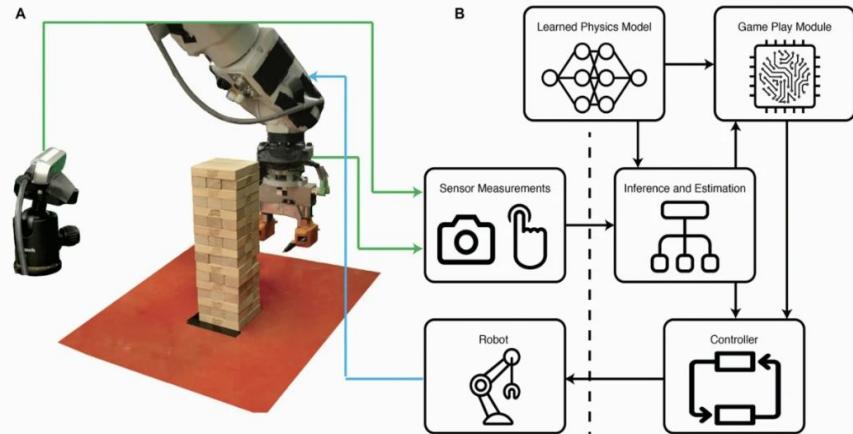
- ICML 2020 Tutorial on
Model Based Reinforcement Learning:

<https://sites.google.com/view/mbrl-tutorial>

Motivation - sample efficiency



Model-based reasoning for ***robotic control***



Fazeli et al. (2019). See, feel, act: Hierarchical learning for complex manipulation skills with multisensory fusion. *Science Robotics*, 4(26).

Why do we want to learn a model?



Planning with real robots
(too expensive, too risky)



**Simulating complex
physical dynamics**
(too expensive)



Interactions with humans
(no access)

Model-free vs model-based

- Collect data $D=\{s_t, a_t, r_{t+1}, s_{t+1}\}$
- Model-free
 - data → policy
- Model -based
 - Data → model → policy

What is a model?

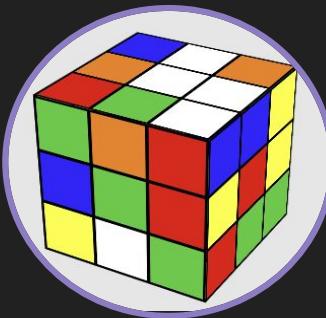
- Some explicit representation of the knowledge about the environment and task
- Usually:
 - Transition/dynamics model: $s_{t+1} = f(s_t, a_t)$
 - Reward model: $r_{t+1} = f(s_t, a_t)$
- Other cases:
 - Inverse transition (going backwards in time)
 - Modeling distance between states

Environment simulators

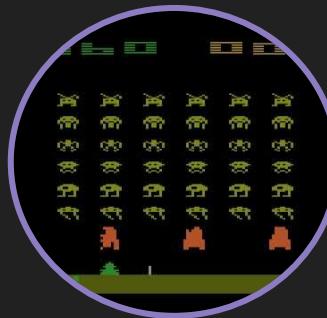
Sometimes we know the ground truth dynamics and rewards.
Might as well use them!



*Silver et al. (2016).
Mastering the game of Go
with deep neural networks
and tree search. Nature.*



*Agostinelli et al. (2019).
Solving the Rubik's cube
with deep reinforcement
learning and search.
Nature Machine
Intelligence.*

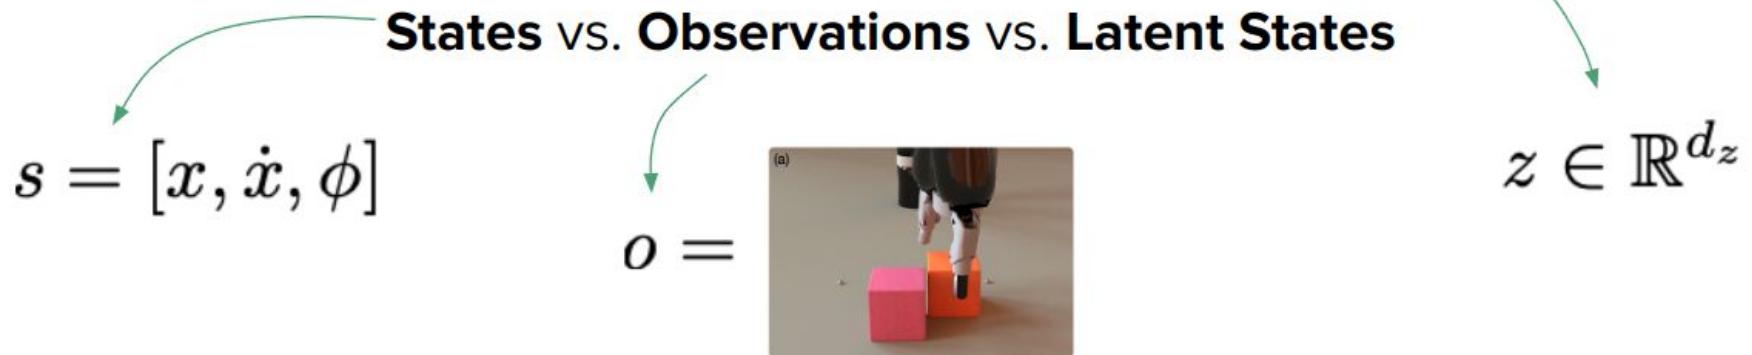


*Bellemare et al. (2013). The
Arcade Learning
Environment: An
Evaluation Platform for
General Agents. JAIR.*



*Todorov et al. (2012).
MuJoCo: A physics engine
for model-based control.
IROS.*

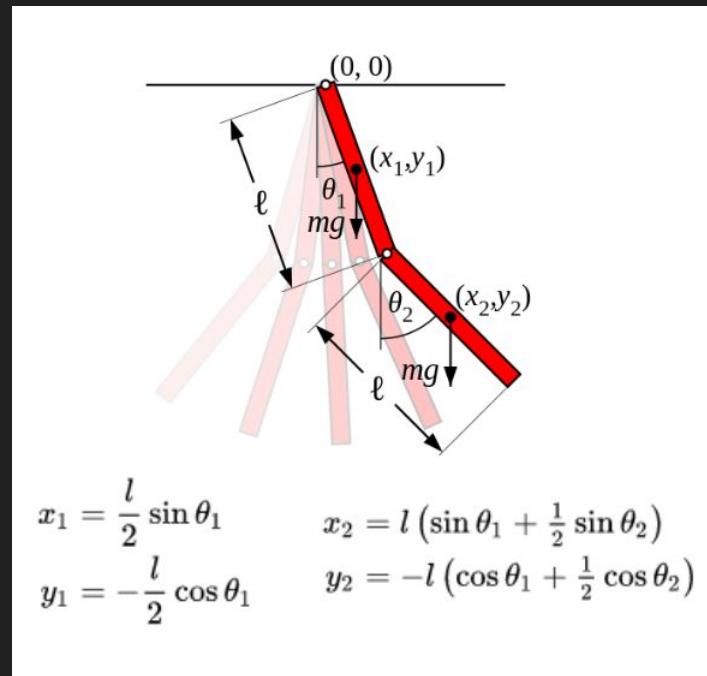
Dimensions of learned models



State transitions: dynamical systems



Double pendulum (acrobot)

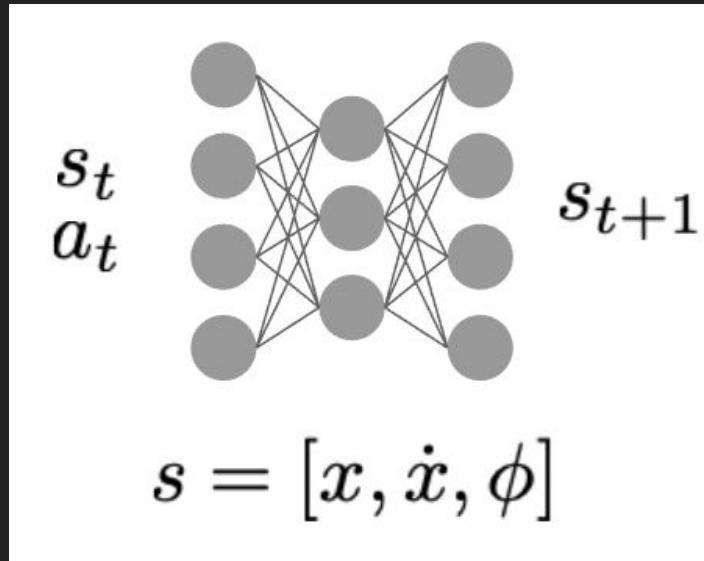


Equations of motion are assumed known.

Reward usually assumed known.

Use system identification to estimate unknown parameters (e.g. mass)

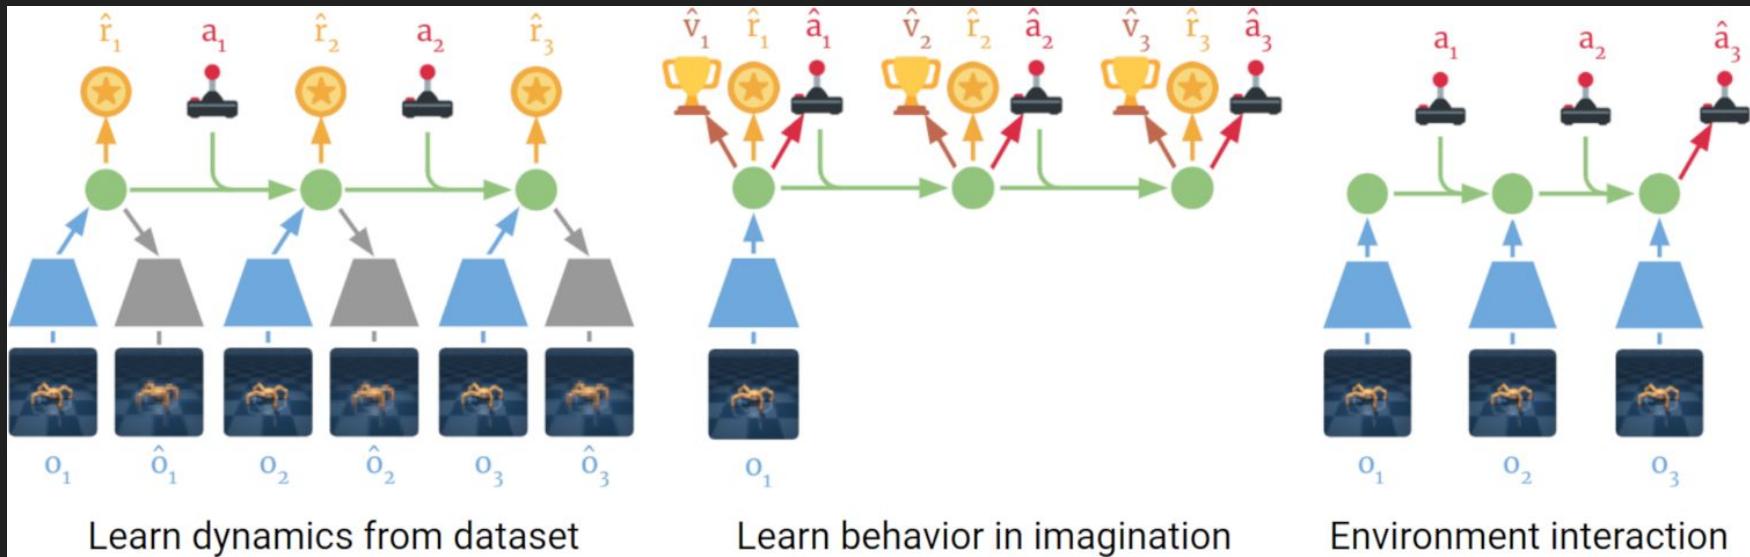
State transitions: MLP-based



Note: typically better to predict the derivative (change in s), and then integrate to obtain s_{t+1}

This is the same idea
behind using skip
connections / residual
blocks!

Latent-space transition models



Learn dynamics from dataset

Learn behavior in imagination

Environment interaction

Model-free vs. model-based RL

Model-free	Model-based	
		Asymptotic rewards*
	/	Computation at deployment*
		Data efficiency*
		Speed to adapt to changing rewards*
		Speed to adapt to changing dynamics*
		Exploration*

*but it depends a lot on the specific method!

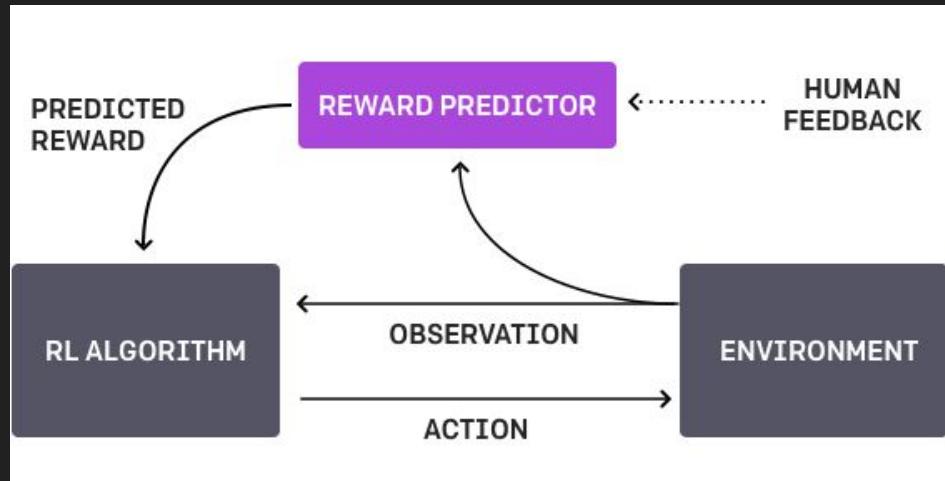
From GPT-3 to ChatGPT

<https://openai.com/blog/instruction-following/>

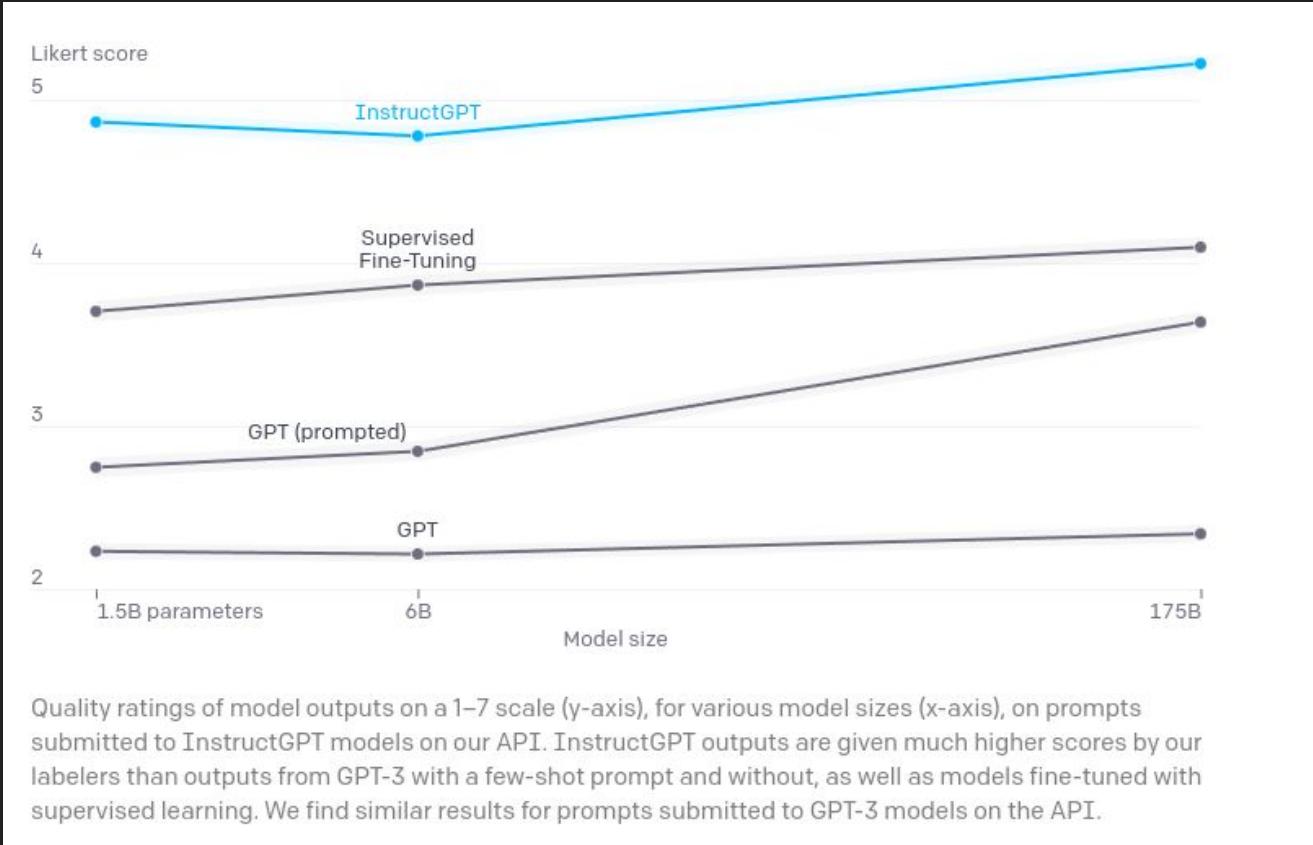
<https://openai.com/blog/chatgpt/>

InstructGPT

- GPT is trained to predict the next token, not to solve a problem posed by a user
- Goal - use human feedback to guide a higher level policy



InstructGPT - results



Methods

Step 1

Collect demonstration data, and train a supervised policy.

A prompt is sampled from our prompt dataset.

Explain the moon landing to a 6 year old

A labeler demonstrates the desired output behavior.



Some people went to the moon...



This data is used to fine-tune GPT-3 with supervised learning.

Step 2

Collect comparison data, and train a reward model.

A prompt and several model outputs are sampled.

Explain the moon landing to a 6 year old

A Explain gravity...
B Explain war...

C Moon is natural satellite of...
D People want to the moon...

A labeler ranks the outputs from best to worst.



D > C > A = B

This data is used to train our reward model.



D > C > A = B

Step 3

Optimize a policy against the reward model using reinforcement learning.

A new prompt is sampled from the dataset.

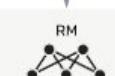
Write a story about frogs

The policy generates an output.



Once upon a time...

The reward model calculates a reward for the output.

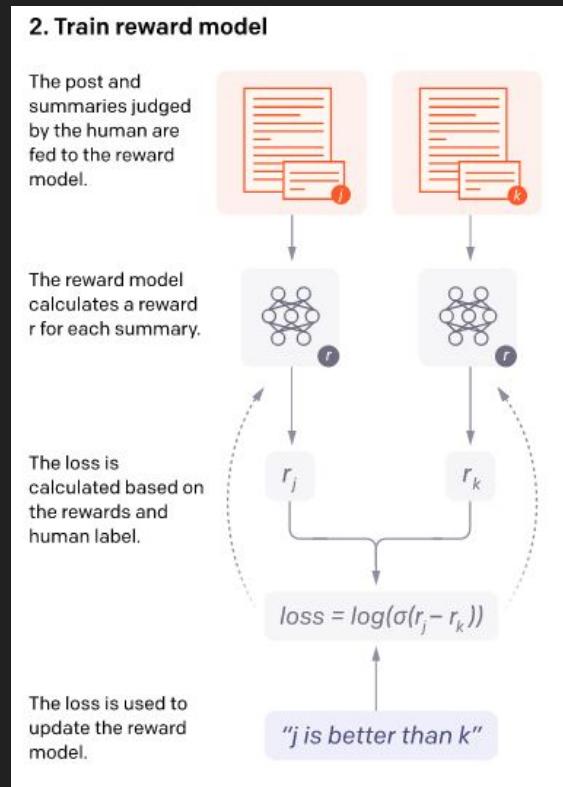


The reward is used to update the policy using PPO.

r_k

Reward model details - work prior to GPT-3

- The reward model calculates absolute reward estimate
- The difference between the absolute values is used to propagate the ranking-based loss.



RL 3-lecture series recap

- Q-learning (DQN)
- Policy Gradients
- Model based RL:
 - Alpha Go Zero
 - Human Feedback RL → InstructGPT (ChatGPT)

Feedback is a gift



Please fill
the course survey in USOS